Vega FEM Library (v4.0) User's Manual*

May 16, 2019

1   Introduction

Vega FEM is a computationally efficient and stable middleware C/C++ library (about 145,000 lines of code) to timestep nonlinear three-dimensional deformable models. It is designed to model large deformations, including geometric and material nonlinearities, and can also efficiently simulate linear elasticity. Vega is a *middleware* physics library; it is not an end-product, or plug-and-play software. Its strength lies in its many C/C++ libraries which depend minimally on each other, and are in most cases independently reusable. Vega is open-source and free, and can be downloaded from `http://www.jernejbarbic.com/vega`. It is released under the 3-clause BSD license, which means that it can be used both in academic research and in commercial applications; see `LICENSE.txt`. This documentation aims to explain how to compile and use Vega. It also provides detail about the code structure to simplify component re-use and the addition of new features.

   Vega is a library for simulation of 3D solids, as well as cloth. The input to Vega is a 3D volumetric mesh with material properties, as well as (time-varying) external forces to act on the 3D mesh at each timestep. Vega computes the displacements of mesh vertices at each timestep. It can also calculate the elastic energy, internal elastic forces and their gradients (tangent stiffness matrix), in any object configuration. Two mesh elements are supported: linear tetrahedra and axis-aligned cubes. Vega includes libraries for sparse matrix storage and arithmetics, volumetric mesh storage (includes a parser for `.veg` format) triangle mesh storage (includes a parser for `.obj` format, and a half-edge data structure implementation), and basic geometric operations, code performance timing, and optional mesh rendering using OpenGL. For sparse linear system solving, Vega can use PARDISO [PAR], SPOOLES [SPO], or its own conjugate gradient implementation, written by following the well-known reference [She94]. Vega also includes an elaborate example driver (executable) simulator to interactively test the implemented materials, timestepping methods, and other settings. Several example meshes and configuration files are included to demonstrate the various materials.

---

*Authors of Vega are Jernej Barbič, Funshing Sin and Daniel Schroeder. This user's manual was written by Jernej Barbič and Daniel Schroeder, and improved by Yili Zhao, Yijing Li, Hongyi Xu, Bohan Wang and Danyong Zhao.
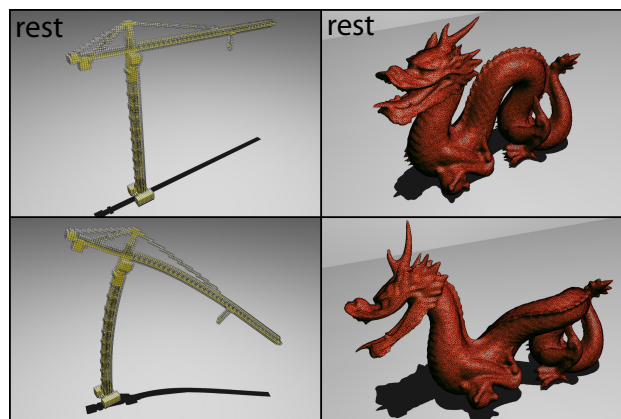

Figure 1: **Large FEM deformations:** The crane cubic (voxel) mesh is shown blended on top of the embedded triangle mesh, whereas the dragon uses the external surface of the tetrahedral mesh for rendering. Original triangle models are courtesy of Waldemar Barkowski, `http://www.artworx-media.de`, and Stanford Computer Graphics Laboratory, `http://graphics.stanford.edu/data/3Dscanrep/`, respectively.

We provide a lot of usage advice and detailed experiments on the performance of Vega in [SSB12]. Please refer to this paper for the theoretical background and explanation of the implemented methods.

Vega provides co-rotational linear Finite Element Method (FEM) elasticity [MG04] (and the exact tangent stiffness matrix [Bar12], similar to [CPSS10]), orthotropic linear materials (combined with co-rotational FEM) [LB14], Saint-Venant Kirchhoff FEM deformable models (see [Bar07]), invertible FEM models [ITF04, TSIF05], and mass-spring systems. All models can simulate quality large deformations and include support for multi-core computing. Vega provides linear materials, as well as neo-Hookean, Saint-Venant Kirchhoff and Mooney-Rivlin nonlinear materials (combined with invertible FEM [ITF04, TSIF05]). Optional compression resistance term [KTY09] is available to improve simulation stability for soft materials. Arbitrary nonlinear materials can be added to Vega. For isotropic hyperelastic materials, this is as easy as defining an energy function and its first and second derivatives. Implemented integrators include implicit backward Euler [BW98], explicit central differences [Wri02], implicit Newmark [Wri02] and symplectic Euler. Other integrators can be added easily. Vega also includes an implementation of the Baraff-Witkin cloth solver [BW98], model reduction [BJ05], and rigid body dynamics [BW01].

Finally, we find it useful to also state what Vega currently does not support. Vega currently supports 3D solid simulation on volumetric meshes, as well as cloth simulations on triangle meshes; it does not support one-dimensional deformable objects (strands). It uses linearized tets and axis-aligned cubes for mesh elements; higher-order elements are not supported. Vega solves linear systems of equations using preconditioned conjugated gradients or direct sparse solvers. It does not incorporate solver acceleration strategies such as multigrid. It does not implement inverse kinematics, or space-time optimization. Vega does not simulate articulated rigid bodies or fluids. It does not support mesh cutting or other run-time mesh modifications; although mesh material properties could be altered at run-time without much difficulty. While it does not incorporate collision detection or contact resolution, it is possible to use any external collision detection library, compute forces externally, and set them as external forces in Vega. A more elaborate contact solver could be written by using Vega's routines to compute internal forces and stiffness matrices in arbitrary mesh configurations.

New in Vega 2.0: improved isotropic materials (greatly improved behavior for soft materials via compression resistance [KTY09]); model reduction (implementation of Jernej Barbič model reduction work [BJ05]; `reducedStVK`, `integratorDense` classes); cloth solver (implementation of Baraff-Witkin's cloth model [BW98]; `clothBW` class); rigid body simulation (implementation Baraff-Witkin SIGGRAPH 2001 course notes [BW01], `rigidBodyDynamics` class); quaternion arithmetics (`quaternion` class); load/save images in PNG, JPEG, TIFF, TGA, PPM formats (wrappers to libpng, libjpeg and libtiff; TGA and PPM are built-in); `imageIO` class); several bug fixes and documentation improvements.

New in Vega 2.1: orthotropic materials (materials that exhibit different stiffnesses in three orthogonal directions, see [LB14]); documentation improvements; binary `.vegb` format; faster loading of volumetric meshes with many regions; support for the clang compiler (Mac); bug fixes.

New in Vega 2.2: minor bug fixes and speed improvements.

New in Vega 3.0: elastic energy functions for all the materials in Vega; improved cloth implementation [BW98]; signed distance field calculation from manifold meshes; or from general "polygon soups" [XB14]; marching cubes with topological guarantees (similar to [LLVT03, Che95]); quality triangular meshing of distance field level sets [ORY05]; Delaunay tetrahedralization of an arbitrary point set in 3D [Bow81] (including efficiently adding new points); constrained Delaunay 3D tet meshing with refinement [Si08, SG11] (experimental feature); calculating the winding number for triangle meshes [JKSH13]; calculating eigenmodes of a 3D mesh under constraints [HX16]; Laplacian matrix of a tet mesh [XLCB15]; improved PARDISO solver wrapper; and bug fixes.

New in Vega 3.1: Positive-definite invertible materials; improved Delaunay mesher and tet mesher; improved code style; bug fixes.

New in Vega 4.0: Virtual tet meshing for nearly self-colliding and colliding triangle meshes (implementation of our published Siggraph 2018 paper); improved (algorithmically) the multi-threading efficiency of elastic force and stiffness matrix computations (via "stencil force models"); replaced pthreads with Intel TBB for multi-threading; support cubic meshes (in addition to tet meshes) in corotationalLinearFEM; provides an as-rigid-as-possible (ARAP) deformation implementation; all meshers now use the significantly faster exact predicates made available by Jonathan Shewchuk (as opposed to wildMagic); new low-level libraries for basic algorithms/data structures and geometry operations; some file rearrangements and renaming.

1.1   Compiling Vega: Core Functionality

This section explains how to build the "core" functionality of Vega, i.e., (unreduced) simulations of 3D solid elasticity. This functionality has been available in Vega since version 1.0 (with some improvements).

Vega is designed to be cross-platform. We have successfully compiled it on Linux, Windows and Mac OS X. We provide makefiles (`make` command) for Linux and Mac OS X. The easiest approach to build Vega is to go to the Vega root folder, and launch `./build`. If all goes well, this should compile the entire Vega core functionality. Details are below.

The first step is to select either `Makefile-headers/Makefile-header.linux`, or `Makefile-headers/Makefile-header.osx`, depending on your OS. By default, linux is selected. The selection is made by creating a soft link `Makefile-headers/Makefile-header` that points to the selected Makefile. To switch to MAC OS X, use:
`cd Makefile-header;ln -sf Makefile-header.osx Makefile-header`.
Alternatively, you can switch using the provided scripts `Makefile-headers/selectLinux` or `Makefile-headers/selectMACOSX`.

To compile the core libraries and the interactive simulator in one step, move to the `utilities/interactiveDeformableSimulator` folder and run `$ make`. To compile an individual library `libraryName`, along with all the libraries it depends on, move to `libraries/libraryName` and run `$ make`. The script `./build` combines `Makefile-header` selection and the compilation into one step.

To clean all the compiled files (for the interactive simulator and for the core libraries), move to `utilities/interactiveDeformableSimulator` and run `$ make deepclean`. To only clean the simulator folder, run `$ make clean`. Similarly, running `$ make deepclean` in folder `libraries/libraryName` will clean the folder for `libraryName` as well as those of all the other `libraries` libraries it depends on, whereas running `$ make clean` will only clean the `libraries/libraryName` folder. The script `./cleanAll` cleans all the libraries and utilities.

Dependencies: Vega has no required external dependencies for the core functionality. It includes all the necessary components for deformable object simulation. If you want to use the optional model reduction functionality (introduced in version 2.0), you will need the dependencies listed in Section 1.3.1. Vega uses its own CG solver as the default solver for implicit integration. This solver can be substituted for external linear system sparse solvers such as PARDISO or SPOOLES. Other solvers can be easily added as well. The provided demonstration driver uses OpenGL for visualization: it depends on the `GL`, `GLU`, `GLUT` OpenGL libraries. It also depends on the GLUI User Interface Library, which is used for the driver's user interface (buttons, edit boxes, etc.). This external library (LGPL license) is included with the Vega distribution (in `libraries/glui`), unmodified as it was downloaded from `http://glui.sourceforge.net/`. The provided Makefiles will automatically compile GLUI to a shared library (to comply with the LGPL license terms), and link it to the driver. On some systems, for the build to work correctly, you may need to figure out how to satisfy the dependencies on OpenGL and/or GLUI. Default settings for the locations of all libraries are provided in the Makefile-header files for Linux and Mac OS X. If needed, custom locations can be set by modifying the `libraries/include/openGL-headers.h` header. For example, to set custom locations and names of the OpenGL libraries, modify the `LIBRARIES`, `STANDARD_LIBS`, `GLUI_DIR`, `GLUI_INCLUDE`, and `GLUI_LIB` variables in `Makefile-header`.

Using PARDISO: Vega contains code to easily use the PARDISO solver from the Intel MKL library. To use PARDISO, you must first install it. To specify the names and locations of PARDISO libraries, modify the `PARDISO_DIR`, `PARDISO_INCLUDE`, and `PARDISO_LIB` variables in the `Makefile-header`. If needed, adjust the

included MKL header files by modifying the `libraries/include/lapack-headers.h` header file. Next, notify Vega that PARDISO is available by defining (uncommenting) the macro `PARDISO_SOLVER_IS_AVAILABLE` in `libraries/sparseSolver/sparseSolverAvailability.h`. Finally, switch the `integrator` library to PARDISO by enabling (defining) the macro `PARDISO` in the following file: `libraries/integrator/integratorSolverSelection.h`. Then, recompile Vega and the driver.

Using SPOOLES: Vega can also use SPOOLES, a public domain sparse solver, freely available at `http://www.netlib.org/linalg/spooles/`. The library is used by the `SPOOLESSolver` and `SPOOLESSolverMT` classes in the `sparseSolver` library. To use SPOOLES, follow the same steps as with PARDISO. To set the location of the SPOOLES library and header files, modify the `SPOOLES_DIR`, `SPOOLES_INCLUDE`, and `SPOOLES_LIB` variables in `Makefile-header`.

Note that the PARDISO and SPOOLES solvers are not necessary to compile our example driver. By default, Vega uses its preconditioned conjugate solver (PCG), available in `libraries/sparseSolver/CGSolver.h`. Note, however, that Intel MKL gives the fastest computation; about 30% faster than SPOOLES and much faster than PCG. PARDISO or SPOOLES are also more stable than PCG, because they can handle system matrices that have negative eigenvalues (indefinite systems). Finally, note that the make files do not monitor the `openGL-headers.h` or `lapack-headers.h` files, so modifications to these may require cleaning and remaking the project, as explained above.

Using Intel TBB: Starting from version 4.0, Vega does not use `pthreads` or OpenMP anymore; but uses Intel Thread Building Blocks (TBB). For example, Vega uses Intel TBB to accelerate the computation of elastic forces and stiffness matrices. Note that by default, the compilation of multi-threading code is disabled, so that Vega can be compiled without Intel TBB. If you want to enable it, uncomment the following line in `Makefile-header`:

```
#TBBFLAG=-DUSE_TBB
```

You also need to add Intel TBB headers and libraries to the following lines in `Makefile-header`:

```
TBB_INCLUDE=<add here>
TBB_LIB=<add here>
```

Then, recompile Vega.

Microsoft Visual Studio: We provide two MSVS solutions (Visual Studio 2015) for the core functionality and all of Vega, respectively. Users can compile the core functionality by opening "vega-basic.sln" in the Vega root folder. In order to build it on Windows, the following external libraries are required: `glui`, `glut` and `glew`. `glui` is shipped with Vega package. For `glut`, a popular implementation that is still actively maintained is `freeglut` (`http://freeglut.sourceforge.net/`).

We provide Visual Studio property sheets (*.props) to configure Vega compilation: libraries/libraries.props, utilities/utilities.props, utilities/debug-libs.props and utilities/release-libs.props. The file `libraries/libraries.props` defines the include directories and c++ preprocessors, and `utilities/utilities.props` stores the common library file paths and library directories. File `utilities/debug-libs.props` configures the Debug libraries, whereas `utilities/release-libs.props` configures the Release libraries. To configure external library path (lib path and include path), you need to first open the property sheet in Visual Studio, and go to the "User Macros" tab. You can then specify external library paths by modifying the macro values. For example, the macro representing the include path for `freeglut` is "GLUT_INCLUDE_DIR"; you can set it by double-clicking on it. Please set the values of required macros ("GLUT_INCLUDE_DIR" in `libraries/libraries.props`, "GLUT_LIBRARIES" in `utilities/debug-libs.props` and `release-libs.props`, "GLEW_INCLUDE_DIR" in `libraries/libraries.props`, "GLEW_LIBRARIES" in `utilities/debug-libs.props` and `release-libs.props`) before building the core functionality of Vega. Therefore, the build steps are as follows:

1. Set required macros. This can be done using the approach described above.

2. Build `glui`. This can be done by opening `libraries/glui/glui-2.35/src/msvc/glui.sln` and building x64 Debug/Release target.

3. Build vega. This can be done by opening vega-basic.sln in the root folder and building everything.

The header files will be copied to `libraries/include`. The built Debug libraries will be in `libraries/lib/Debug` and the built utilities will be in `utilities/bin/Debug` (and equivalently for Release).

In order to build full Vega functionality, you need to have Windows versions of all required libraries mentioned in this manual, and you need to specify the correct "User Macros" values for them (same procedure as in the previous paragraph). After properly configuring all necessary libraries, you can open vega.sln and build everything. Note that for libraries `sparseSolver` and `sceneObjectReduced` you need to add all source files to the project before building full Vega functionality. In addition, to enable the multi-threading feature, you need to set the correct Intel TBB paths for macros in all .props files and add preprocessor "USE_TBB" to `libraries/libraries.props` $\rightarrow C/C++ \rightarrow Preprocessor \rightarrow PreprocessorDefinitions$.

For libraries that are not included in the solution, you can create a new project as follows. First, create a new win32 console project in the Vega solution and choose the static library option. Then, add all the source files into the project, and set the target to x64 (our solution is not set up for win32). Finally, add `libraries/libraries.props` and `libraries/copy-headers.props` to the project. This can be done in the Property Manager in Visual Studio. For utilities that are not included in the solution, you can create a new project as follows. First, create a new win32 console project in the Vega solution. Then, add all the source files into the project, and set the target to x64. Finally, add `utilities/utilities.props` and add `utilities/debug-libs.props` and `utilities/release-libs.props` to the project for Debug and Release configuration, respectively.

Because all the Vega projects are configured as Multi-Thread DLL (default setting in Visual Studio), the executables require the DLLs for the dependent external libraries. Therefore, you need to copy those DLLs either into a system folder, or into a folder with the executables; otherwise, you will get a runtime error. After you are able to launch the executables, you can test the Vega examples either by using the Windows command line, or using Visual Studio. Linux scripts (run_***) do not work on Windows, but you can use them as a reference.

1.2 Driver Executable and Examples

Vega provides a complete simulation driver (executable), as well as several complete examples. In order to build the executable driver, follow the instructions from the previous sections. The provided examples include all the necessary meshes and configuration files. To run the simulator on these configurations, run any of the `run_[modelname]` shell scripts in the `examples` folder. For example, try `examples/turtle/run_turtle`, `examples/asianDragon/run_asianDragon` or `examples/asianDragon/run_asianDragon_unconstrained`. You can read the `examples/guideToExamples.txt` for a brief overview of our examples. The simulator executable `interactiveDeformableSimulator` takes one command-line argument: a `.config` configuration file such as those in the `examples` folder. The `.config` file specifies such information as the mesh, the desired material and timestepping method to use, the timestep, damping, an optional secondary rendering mesh embedded in the deformable mesh, and so on. The simulator must be run in the folder containing the desired `.config` file, since this file contains relative paths to other files read by the executable. It is easiest to achieve this by using the provided `run*` shell scripts. Inside the simulator, some useful keys are 'E', 'e', and 'w', to toggle the displays of embedded rendering mesh, volumetric mesh, and volumetric mesh wireframe, respectively. Note that the main simulation window **must have focus** for the keys to work.

1.3 Compiling Vega: Model Reduction

In order to build model reduction, you must first build the "core" functionality (previous section).
Important: For model reduction, PARDISO solver must be available. Therefore, you must install it, and uncomment the `#DEFINE PARDISO_SOLVER_IS_AVAILABLE` line in
`libraries/sparseSolver/sparseSolverAvailability.h`. This step must be done ****before**** building the core Vega functionality. In case you already previously compiled "core" functionality, go to

`utilities/interactiveDeformableSimulator` and run `make deepclean`, then fix `sparseSolverAvailability.h` and recompile "core" functionality using `./build`.

Before you can build model reduction, you must also satisfy all the dependencies (below). Once you have done so, you can issue the `./buildModelReduction` script in the main Vega folder. If all goes well, this should compile all the model reduction libraries, as well as the demonstration driver `reducedDynamicSolver-rt` and the preprocessor application `LargeModalDeformationFactory`.

1.3.1 Model Reduction Dependencies

The model reduction part of Vega has the following dependencies: Intel MKL (PARDISO), GLEW, ARPACK, wxWidgets, and optionally, Cg (Nvidia). You must install all the required dependencies and edit the `Makefile-headers/Makefile-header` file accordingly.

- **Intel MKL**: provides a dense and sparse linear solver. Both a dense and sparse solver are required for model reduction. On some platforms, such as Mac OS X, the dense solver is built-in into the operating system (the framework "Accelerate"), and therefore Intel MKL is not needed for the dense solves. For the sparse linear solver, we typically use the PARDISO solver from Intel MKL, but SPOOLES is a good free alternative. By using SPOOLES or PCG, and an external dense solver, it is possible to completely eliminate the need for Intel MKL. To control what solver is used, set the macros in `libraries/sparseSolver/sparseSolverAvailability.h`.

- **GLEW**: needed to access OpenGL extensions. This is a standard, free, library.

- **Cg**: needed for GP-GPU shaders to compute vertex deformations ($u = Uq$) on the GPU (`ObjMeshGPUDeformer` library). This is a free library that can be downloaded from the Nvidia website. If you don't add Cg libraries in `Makefile-header`, Vega will use CPU code to compute vertex deformations.

- **ARPACK**: needed to solve large sparse eigenproblems in the `LargeModalDeformationFactory`. Available free of charge (BSD license) at `http://www.caam.rice.edu/software/ARPACK/`. In order to compile ARPACK, you will need a Fortran compiler. We use `gfortran`, the standard Fortran compiler that ships with `gcc` on Unix systems.

- **wxWidgets**: this is the UI library used in the `LargeModalDeformationFactory`. Available freely at `http://www.wxwidgets.org`. Note: In newer versions of Mac OS X, we were only able to compile this library with clang, but not with gcc.

1.3.2 Model Reduction Drivers

LargeModalDeformationFactory: Given a simulation mesh (tetrahedral or voxel), its material properties (may be non-homogeneous), and a set of constrained vertices, this application driver generates a fast reduced deformable model suitable for fast large-deformation simulation, as described in [BJ05]. The driver supports the St. Venant-Kirchhoff material, i.e., linear isotropic material parameterized by Young's moduli and Poisson's ratios. The application starts either with a triangle mesh or a volumetric mesh, and generates all the necessary data for a subsequent real-time simulation. You can then run a real-time simulation from your own code using `ImplicitNewmarkDense` and `StVKReducedInternalForces` classes. The application supports tetrahedral and voxel 3D volumetric meshes. The application also contains several generally useful sub-components: compute linear modes (via ARPACK) of arbitrary volumetric (tet or voxel) meshes, under arbitrary boundary conditions (including free boundary conditions), compute modal derivatives, compute volumetric mesh mass and stiffness matrix, create cube volumetric meshes for arbitrary input triangle geometry. The application outputs several files, which are loadable by various Vega classes. For example, it produces the cubic polynomial of reduced internal forces that you can then load using the `StVKReducedInternalForces` and `ImplicitNewmarkDense` classes and timestep in your own code. All the modal matrices generated (linear modes, modal derivatives, or simulation basis matrix) can be loaded using the "Matrix" class. If you want to use tetrahedral meshes, you can generate them using the TetGen mesh generation package or our experimental tet mesher introduced in Version 3.0.

reducedDynamicSolver-rt: This is an example driver that demonstrates how to launch a run-time model reduction simulation (created via `LargeModalDeformationFactory`). In `utilities/reducedDynamicSolver-rt`, there are two example demos: `run_A`, `run_simpleBridge`.

1.4   Compiling Vega: Tet Mesher

Since Version 3.0, Vega provides the ability to compute distance fields, create iso-surface meshes (via marching cubes), and create tet meshes using constrained Delaunay tetrahedralization. There is no external dependency for these functionalities. The script `./buildTetMeshers` in the main Vega folder compiles the corresponding drivers. Examples can be found in `examples/mesher`.

1.5   Compiling Vega: Immersion Mesher

Since version 4.0, Vega includes a tet mesher for nearly self-intersecting, and self-intersecting geometry. In order to build these libraries, as well as their demo drivers, you need additional dependencies: CGAL, Eigen and libigl. PARDISO solver (such as the one included in Intel MKL) is also recommended, as it needed in the as-rigid-as-possible [SA07] solver used for visualizing the results.

- **CGAL**: The CGAL's exact arithmetic kernel (licensed under LGPL) is needed for nearly self-intersecting geometry, whereas self-intersecting geometry requires the GPL part of CGAL.

- **libigl**: It is needed to compute the cell complex for our immersion mesher. This functionality of libigl depends on Eigen and CGAL. Both libigl and Eigen are licensed under MPL2.

- **Eigen**: Free linear algebra library used by libigl.

Once the dependencies are added in `Makefile-header`, run the script `./buildImmersionMesher` to compile the drivers. Examples can be found in `examples/virtualTets` and `examples/immersionMesher`.

1.6   Short Tutorial on Using Vega

Vega is a middleware library and should be integrated with the rest of the user's C/C++ code. In this section, we give the typical steps to do so. Details on the specific methods can be found in subsequent sections. The first step is to load a volumetric mesh from a `.veg` file (format is described in Section 1.7), which creates a `VolumetricMesh` object. It is easiest to do so using the provided `VolumetricMeshLoader` class, which automatically detects the type of elements in the mesh (tets or cubes).

```
#include "volumetricMeshLoader.h"
...

char inputFilename[96] = "myInputMeshFile.veg";
VolumetricMesh * volumetricMesh = VolumetricMeshLoader::load(inputFilename);
if (volumetricMesh == NULL)
  printf("Error: failed to load mesh.\n");
else
  printf("Success. Number of vertices: %d . Number of elements: %d .\n",
    volumetricMesh->getNumVertices(), volumetricMesh->getNumElements());
```

At this point, the mesh and its material properties have been parsed into memory, and we have a valid `VolumetricMesh` object. Next, we will initialize a specific 3D deformable model, allowing us to compute internal forces and stiffness matrices for arbitrary deformed object configurations. Let us use the corotational linear FEM model.

```
#include "corotationalLinearFEM.h"
...

CorotationalLinearFEM * deformableModel = new CorotationalLinearFEM(volumetricMesh);
```

We now have a valid deformable model, and it is possible to query internal forces and tangent stiffness matrices in any mesh configuration. Typically, however, we want to timestep the model in time. So let's proceed with building an integrator for the model. We need to first create a `stencilForceModel` and `ForceModel` object, to connect our deformable model to an integrator. The usage of the `stencilForceModel` is needed to enable the fast multi-core functionality available since Vega FEM v4.0. We also need the mass matrix, and specify which model vertices (if any) are to be fixed. Then, we can initialize the integrator. Let us use the implicit backward Euler integrator. Note that the sparse linear system solver to use for implicit integration is selected by editing the file "integratorSolverSelection.h" inside the `Integrator` library. The default selection is to use Vega's conjugate gradient solver.

```
#include "corotationalLinearFEMStencilForceModel.h"
#include "forceModelAssembler.h"
#include "generateMassMatrix.h"
#include "implicitBackwardEulerSparse.h"
...

// Create the class to compute internal elastic forces and stiffness matrices
// on each mesh element.
StencilForceModel * stencilForceModel =
  new CorotationalLinearFEMStencilForceModel(deformableModel);

// Create the class to connect the deformable model to the integrator.
// ForceModelAssembler automatically computes in parallel if Intel TBB is enabled.
ForceModel * forceModel = new ForceModelAssembler(stencilForceModel);

int r = 3 * tetMesh->getNumVertices(); // total number of DOFs

double timestep = 0.0333; // the timestep, in seconds

SparseMatrix * massMatrix;
// Create consistent (non-lumped) mass matrix.
GenerateMassMatrix::computeMassMatrix(tetMesh, &massMatrix, true);

// constraining vertices 4, 10, 14 (constrained DOFs are specified 0-indexed):
int numConstrainedDOFs = 9;
int constrainedDOFs[9] = { 12, 13, 14, 30, 31, 32, 42, 43, 44 };

// (tangential) Rayleigh damping
double dampingMassCoef = 0.0; // "underwater"-like damping (here turned off)
double dampingStiffnessCoef = 0.01; // (primarily) high-frequency damping

// initialize the integrator
ImplicitBackwardEulerSparse * implicitBackwardEulerSparse = new
  ImplicitBackwardEulerSparse(r, timestep, massMatrix, forceModel,
    numConstrainedDOFs, constrainedDOFs,
    dampingMassCoef, dampingStiffnessCoef);
```

At this point, we can start timestepping our model! By default, initial conditions are zero deformation and zero velocity. Arbitrary initial conditions could be set via `IntegratorBase::SetState`. Let us apply some forces to the model, perform a couple of timesteps, and read the results. The forces are specified in Newtons (N).

```
// alocate buffer for external forces
double * f = (double*) malloc (sizeof(double) * r);
int numTimesteps = 10;
```

```
for(int i=0; i<numTimesteps; i++)
{
  // important: must always clear forces, as they remain in effect unless changed
  implicitBackwardEulerSparse->SetExternalForcesToZero();
  if (i == 0) // set some force at the first timestep
  {
    for(int j=0; j<r; j++)
      f[j] = 0; // clear to 0
    f[37] = -500; // apply force of -500 N to vertex 12, in y-direction, 3*12+1 = 37
    implicitBackwardEulerSparse->SetExternalForces(f);
  }
  implicitBackwardEulerSparse->DoTimestep();
}


// alocate buffer to read the resulting displacements
double * u = (double*) malloc (sizeof(double) * r);
implicitBackwardEulerSparse->GetqState(u);
```

And that's it – these were all the necessary steps to use Vega! The array u now contains the mesh vertex displacements after 10 simulation timesteps. We can now continue simulating, render the object, or perform collision detection (using some external software). We could then set the resulting contact forces as the external forces so that they will affect the deformations in the subsequent steps.

Choosing the timestep: The unit for the timestep is seconds (s). It is very important to choose a sufficiently small timestep so that the simulation is stable. Too large timesteps will lead to simulation blow-up. If the simulation is unstable, the first step should always be to greatly decrease the timestep and see if the simulation is stable. Some integrators are inherently more stable than others, e.g., implicit backward Euler is more stable than implicit Newmark (but introduces more artificial damping). Also, direct sparse solvers (PARDISO and SPOOLES) tend to be more stable than Conjugate Gradients. For many more such usage tips, refer to our publication [SSB12].

1.7   Vega file formats

In this section, we document the Vega file format, .veg. This is a text (ASCII) file format which is flexible and the files are easy to edit manually; we recommend that new users familiarize themselves with this format first. Since Vega FEM 2.1, Vega also supports a binary file format, .vegb, which can be used when smaller file sizes and faster loading times are desired. The usage of .vegb is documented in the VolumetricMesh, TetMesh and CubicMesh classes.

The .veg format extends the open source volumetric mesh file format developed by Jonathan Shewchuk and employed by Stellar [KS09, Kli08] and TetGen [Han11] mesh generation packages. Meshes in Shewchuk's format are easily loadable into Vega. Note that the Stellar webpage [KS09] contains a converter script that can convert other popular mesh formats into the Stellar/TetGen format.

For simulation, it is necessary to specify mesh material properties, in addition to geometry. The Vega file format .veg achieves this by introducing a set of additional keywords for material specification. Vega supports *heterogeneous material properties*, i.e., different parts of the mesh can have different material properties. In the extreme case, every mesh element can have a separate set of material properties. We provide several example .veg meshes in the models folder. For heterogeneous material properties, see the turtle example in models/turtle.

Basics: The Vega file format is ASCII. Lines starting with * denote a command. Lines starting with # are comments. Empty lines are ignored. Files can be nested using the *INCLUDE command. The effect of *INCLUDE is to include, at that point in the .veg file, the entire contents of the included file. Include files can include other files; they can nest arbitrarily.

Vertices are specified using the `*VERTICES` keyword. The second line gives the number of vertices, followed by integer "3" (three-dimensional simulation), optionally followed by more parameters, which are ignored. The subsequent lines give one vertex per line, in the format:

`<vertex index> <x> <y> <z>`

where the vertex index starts either at 0 or 1, increments by 1 for every vertex. The unit for vertex positions is meters [m]. Example:

```
*VERTICES
5 3 0 0
1 0.5 0.5 0.5
2 1.0 -0.5 0.5
3 -1.0 0.0 1.0
4 0.25 -0.25 0.5
5 0.6 0.2 0.3
```

Elements are specified using the `*ELEMENTS` keyword. The second line gives the element type, either "TET" or "CUBIC". Tetrahedra can have arbitrary shape. Of course, degenerate tetrahedra should be avoided, whereas tetrahedra with small dihedral angles generally cause a higher linear system condition number, which may cause instabilities. Use of a quality tet mesh generator is advisable. Cubic meshes must consist of cubes: all sides of cubes must be equal; general cuboids (boxes) are not supported. Cubes must be axis-aligned. The third line gives the number of mesh elements, followed by the number of vertices in each element (4 for tets and 8 for cubes), followed by some optional integers that are ignored. Subsequent lines give one element per line, in the format:

`<element index> <vertex 1> ... <vertex n>`

where the element index starts either at 0 or 1 and increments by 1 for every element, and n is the number of element vertices. For tet, $n = 4$, for cubes, $n = 8$. Example:

```
*ELEMENTS
TETS
2 4 0
1 2 3 4 1
2 5 3 4 1
```

specifies a mesh with two tets. The first tet has vertices $2, 3, 4, 1$ (in that order), and the second tet has vertices $5, 3, 4, 1$. The ordering of vertices within the element matters. Incorrect order will give wrong results. For tet meshes, tets must be oriented so that vertices $1, 2, 3, 4$ specify the tet in a positive orientation, i.e., $((v_2 - v_1) \times (v_3 - v_1)) \cdot (v_4 - v_1) \geq 0$. If the input mesh does not have this property (or you are unsure), you can call the function `orient` in the `TetMesh` class to establish a correct orientation. Note that the meshes produced by TetMesh have positive orientation. For cubic meshes, the vertices must be specified in the following order: $(0, 0, 0), (1, 0, 0), (1, 1, 0), (0, 1, 0), (0, 0, 1), (1, 0, 1), (1, 1, 1), (0, 1, 1)$. This example specifies the unit cube; cubes whose lower-left-front corner is not at the origin and/or whose size is not unit must be specified analogously. Other orders will give wrong results.

New in Vega 2.0: vertices and elements can now be numbered starting either at 0, or at 1. They must be given in a sorted, ascending order, starting either from 0 or 1. Previous versions of Vega only supported starting at 1. Support for 0 was added to accommodate the recent versions of TetGen, which produce tet meshes with vertices and elements starting at 0. The following tet mesh is valid in Vega 2.0, and equivalent to the tet mesh given above:

```
*VERTICES
5 3 0 0
0 0.5 0.5 0.5
1 1.0 -0.5 0.5
```

```
2 -1.0 0.0 1.0
3 0.25 -0.25 0.5
4 0.6 0.2 0.3
*ELEMENTS
TETS
2 4 0
0 1 2 3 0
1 4 2 3 0
```

Materials are specified using the *MATERIAL keyword. The first line gives the material name, the second line specifies the material properties. Three material specifications are supported: "ENU", for any material that is parameterized by Young's modulus (E, in N/m^2) and Poisson's ratio (ν, dimensionless quantity (no unit)), "MOONEYRIVLIN" for Mooney-Rivlin materials, and "ORTHOTROPIC" for orthotropic materials. All materials include a mass density specification, in kg/m^3.

Most materials in Vega use the "ENU" specification: co-rotational linear FEM, StVK, invertible StVK, invertible neo-Hookean, etc. Example:

```
*MATERIAL mat1
ENU, 1000, 1E8, 0.40
```

specifies a material with mass density $1000\,kg/m^3$, Young's modulus of $10^8\,N/m^2$, and Poisson's ratio of 0.4.

The Mooney-Rivlin material can be simulated using the invertible FEM class `IsotropicHyperelasticFEM` and is specified as follows:

```
*MATERIAL myMaterialName
MOONEYRIVLIN, 500, 0.5, 0.6, 1.0
```

specifies a Mooney-Rivlin material with mass density $500\,kg/m^3$, $\mu_{01} = 0.5, \mu_{10} = 0.6$ and $v_1 = 1.0$. The implemented Mooney-Rivlin material is described in Section 3.5.5 of [Bow09].

Orthotropic materials [LB14] are materials that exhibit different stiffnesses in three orthogonal directions. By default, the orthogonal directions are $(1, 0, 0), (0, 1, 0), (0, 0, 1)$, but arbitrary orthogonal directions are supported. The implemented orthotropic material supports large deformations and is described in [LB14]. Orthotropic materials in Vega can be simulated using the corotational linear FEM deformable model available in the `CorotationalLinearFEM` class. Vega supports several formats to specify an orthotropic material, based on whether the user wants to specify an orthotropic material in full generality, or how many values the user wants to leave at the default settings.

```
*MATERIAL myMaterialName
ORTHOTROPIC, 500, 1E8, 2E8, 3E8, 0.4, 0.45, 0.5, 5E7, 8E7, 7E7, 0.866025, -0.5,
0, 0.5, 0.866025, 0, 0, 0, 1
(the entire specification starting from ORTHOTROPIC must be on one line)
```

This is the default format for orthotropic materials. It specifies the parameters for an orthotropic material with mass density $500\,kg/m^3$, Young's moduli ("stiffnesses") $E_1 = 10^8\,N/m^2, E_2 = 2 \times 10^8\,N/m^2, E_3 = 3 \times 10^8\,N/m^2$ in the three orthogonal directions, Poisson's ratios ("volume preservation coefficient") $\nu_{12} = 0.4, \nu_{23} = 0.45, \nu_{31} = 0.5$, shear moduli $\mu_{12} = 5 \times 10^7\,N/m^2, \mu_{23} = 8 \times 10^7\,N/m^2, \mu_{31} = 7 \times 10^7\,N/m^2$ and a 3×3 rotation matrix Q in row-major order representing the orientation of the three orthotropic principal axes. Note that there are six Poisson's ratios ν_{ij}, for $i \neq j$, only three of which are independent. Formulas for $\nu_{21}, \nu_{32}, \nu_{13}$ are in [LB14]. Briefly, E_i gives the stiffness of the material when loaded in orthogonal direction i. Poisson's ratio ν_{ij} gives the contraction in direction j when the material extends (stretches) in direction i. See [LB14] for details on the meaning of the parameters. Young's moduli must satisfy the conditions $E_i > 0$, for $i = 1, 2, 3$. This format gives the user full control over orthotropic materials. However, the user

must ensure that the elasticity tensor of this material is positive-definite, otherwise the simulation might be unstable. The conditions involve all E_i and ν_{ij}, and are available in [LB14]. The rotation Q is

$$Q = \begin{bmatrix} 0.866025 & -0.5 & 0 \\ 0.5 & 0.866025 & 0 \\ 0 & 0 & 1 \end{bmatrix}, \tag{1}$$

i.e., a rotation by 30 degrees around the z-axis. First orthogonal axis is $(0.866025, 0.5, 0)$ (stiffness $E_1 = 10^8 \, N/m^2$), second axis is $(-0.5, 0.866025, 0)$ (stiffness $E_2 = 2 \times 10^8 \, N/m^2$), and the third axis is $(0, 0, 1)$ (stiffness $E_3 = 3 \times 10^8 \, N/m^2$) in this example.

```
*MATERIAL myMaterialName
ORTHOTROPIC_N3G3R9, 500, 1E8, 2E8, 3E8, 0.4, 0.45, 0.5, 5E7, 8E7, 7E7, 0.866025, -0.5,
0, 0.5, 0.866025, 0, 0, 0, 1
(the entire specification starting from ORTHOTROPIC_N3G3R9 must be on one line)
```

specifies the same material as above with all independent parameters. In other words, ORTHOTROPIC and ORTHOTROPIC_N3G3R9 are identical.

```
*MATERIAL myMaterialName
ORTHOTROPIC_N3G3, 500, 1E8, 2E8, 3E8, 0.4, 0.45, 0.5, 5E7, 8E7
```

specifies an orthotropic material with mass density $500 \, kg/m^3$, Young's moduli $E_1 = 10^8 \, N/m^2, E_2 = 2 \times 10^8 \, N/m^2, E_3 = 3 \times 10^8 \, N/m^2$, Poisson's ratios $\nu_{12} = 0.4, \nu_{23} = 0.45, \nu_{31} = 0.5$, shear moduli $\mu_{12} = 5 \times 10^7 \, N/m^2, \mu_{23} = 8 \times 10^7 \, N/m^2, \mu_{31} = 7 \times 10^7 \, N/m^2$ and a default orientation where the three principal axes are aligned with the world coordinate axes, i.e., Q is identity.

```
*MATERIAL myMaterialName
ORTHOTROPIC_N1G1R9, 500, 1E8, 2E8, 3E8, 0.4, 1.1, 0.866025, -0.5,
0, 0.5, 0.866025, 0, 0, 0, 1
(the entire specification starting from ORTHOTROPIC_N1G1R9 must be on one line)
```

specifies an orthotropic material in a simplified, but stable way, using a single Poisson's ratio-like parameter ν, as described in [LB14]. Mass density is $500 \, kg/m^3$, Young's moduli are $E_1 = 10^8 \, N/m^2, E_2 = 2 \times 10^8 \, N/m^2, E_3 = 3 \times 10^8 \, N/m^2$, Poisson's ratio-like parameter is $\nu = 0.4$, and the shear moduli scaling factor is $\mu = 1.1$. This format also specifies a 3×3 rotation matrix Q in row-major order representing the orientation of the three orthotropic principal axes. This format uses the method described in [LB14] to produce a material that is guaranteed to be stable. Parameter ν must satisfy $-1 < \nu < 1/2$. It is used to create a set of stable Poisson's ratios. The user can tune ν like the Poisson's ratio in an isotropic material. This format also uses an automatic method to compute shear moduli from E_1, E_2, E_3 and ν; the formulas are given in [LB14]. Parameter μ scales the shear moduli computed by the method of [LB14]. A value of $\mu = 1$ means that the shear moduli computed by [LB14] will be used unmodified. Values $\mu > 1$ will increase shear stresses (make the model shear less), whereas values $0 < \mu < 1$ will decrease shear stresses (make the model shear more).

```
*MATERIAL myMaterialName
ORTHOTROPIC_N1G1, 500, 1E8, 2E8, 3E8, 0.4, 1.1
```

specifies an orthotropic material in the same way as ORTHOTROPIC_N1G1R9, except that the rotation is assumed to be identity, i.e., a default orientation where the three principal axes are aligned with the world coordinate axes. Mass density is $500 \, kg/m^3$, Young's moduli are $E_1 = 10^8 \, N/m^2, E_2 = 2 \times 10^8 \, N/m^2, E_3 = 3 \times 10^8 \, N/m^2$, Poisson's ratio-like parameter is $\nu = 0.4$, and the shear moduli scaling factor is $\mu = 1.1$.

```
*MATERIAL myMaterialName
ORTHOTROPIC_N1R9, 500, 1E8, 2E8, 3E8, 0.4, 0.866025, -0.5,
0, 0.5, 0.866025, 0, 0, 0, 1
```

specifies an orthotropic material in the same way as ORTHOTROPIC_N1G1R9, except that μ is assigned the default value $\mu = 1$. Mass density is $500\,kg/m^3$, Young's moduli are $E_1 = 10^8\,N/m^2, E_2 = 2{\times}10^8\,N/m^2, E_3 = 3 \times 10^8\,N/m^2$, Poisson's ratio-like parameter is $\nu = 0.4$.

```
*MATERIAL myMaterialName
ORTHOTROPIC_N1, 500, 1E8, 2E8, 3E8, 0.4
```

specifies an orthotropic material in the same way as ORTHOTROPIC_N1G1, except that μ is assigned the default value $\mu = 1$. Mass density is $500\,kg/m^3$, Young's moduli are $E_1 = 10^8\,N/m^2, E_2 = 2{\times}10^8\,N/m^2, E_3 = 3 \times 10^8\,N/m^2$, Poisson's ratio-like parameter is $\nu = 0.4$.

Sets store a set of integer indices. They are used to store indices of elements that share the same material (=region). First line specifies the set name, followed by a comma-separated list of set elements. Elements should be sorted and identified by the indices given in the *ELEMENTS section, i.e., either 0-indexed or 1-indexed, depending on whether *ELEMENTS are given 0-indexed or 1-indexed. Example:

```
*SET set1
11, 17, 21, 37, 113, 220, 310, 555,
556, 557, 570, 601, 991, 1013, 1210, 1225
```

Regions: Elements that share material properties are organized into *regions*. A region is specified using a *REGION keyword. For example,

```
*REGION
set1, material1
```

creates a region consisting of the elements specified in the Set `set1`, and assigns material `material1` to it. In order to set the entire mesh to a material, you can use the built-in set `allElements`:

```
*REGION
allElements, material1
```

If the union of all specified regions does not contain all mesh elements, the remaining elements are assigned a material as follows. The assigned material is the last material mentioned in the `.veg` if the file specified at least one material. If no material was specified, the default material is assigned to the entire mesh. The default material is of type "ENU", with default parameters $E = 10^6 N/m^2$, $\nu = 0.45$, $\rho = 1000 kg/m^3$, where E is Young's modulus, ν is Poisson's ratio and ρ is mass density. If specification of regions was omitted from a `.veg` file, the entire mesh is assigned the default material.

Easy re-use of Stellar/TetGen meshes: Suppose that the mesh vertices and elements are stored in `myMesh.node` and `myMesh.ele` files, in the standard Stellar/TetGen format. The following "template" `.veg` file is the shortest way to import those meshes into Vega:

```
*VERTICES
*INCLUDE myMesh.node
*ELEMENTS
TETS
*INCLUDE myMesh.ele
```

Of course, material properties and regions could be appended as described in the previous paragraphs.

1.8 Acknowledging

If you use Vega, we will appreciate if you acknowledge it. Please use the following citation:

```
Jernej Barbič, Fun Shing Sin, Daniel Schroeder:
Vega FEM Library.  2012.  http://www.jernejbarbic.com/vega
```

Here is the BibTeX file:

```
@misc{Vega,
  author =       "Jernej Barbi\v{c} and Fun Shing Sin and Daniel Schroeder",
  title =        "{Vega FEM Library}",
  year = "2012",
  note = "http://www.jernejbarbic.com/vega",
}
```

1.9 FAQ

1.9.1 How can I create volumetric meshes for use with Vega?

You can use the routines in the `mesher` library. These routines can generate quality tet meshes that mesh an isosurface of a signed distance field. You can calculate the signed distance field using the "signed distance field from polygon soup" pipeline [XB14] available in the `distanceField` library. This pipeline was used, for example, to generate the tet mesh for our turtle, asian dragon and simple bridge (tet) examples. Voxel meshes can be generated using the `objMesh` library, see the `ObjMeshOffsetVoxels` class; this was used for the simple bridge (vox) and tower crane examples. If you need to embed a nearly-self-intersecting manifold surface mesh into a tet mesh, you can use the routines in `virtualTets` library to avoid unnecessarily high tet mesh resolutions around nearly-self-intersecting regions. If your manifold surface mesh is self-intersecting, you can use the `immersionMesher` library to create a tet mesh that conforms to the surface and does not glue the self-intersecting parts. In addition, any 3D tet or cubic mesh can be loaded into Vega (`.veg` file format). You can use any external mesher, such as, for example, TetGen or Stellar. There are many commercial meshers available.

1.9.2 What is the `.veg` file format?

Before Vega, there was no free file format to specify 3D volumetric meshes with material properties. So, we extended the popular free format of Jonathan Shewchuk, which can specify 3D geometry, to also support mesh material properties. Meshes given in Shewchuk's format (such as those produced by TetGen, or Stellar) are trivially loadable into Vega. If material specification is omitted, default material parameters are assigned to the entire mesh. See the above sections for the documentation of the `.veg` file format. The `.veg` file format is free.

1.9.3 What is the `.vegb` file format?

This is a binary file format, offering the same functionality as `.veg`. It was introduced in Vega FEM 2.1, so that files can be made smaller and loading times decreased. So, since Vega FEM 2.1., users have a choice between `.veg` and `.vegb`.

1.9.4 Can Vega simulate non-homogeneous material properties?

Yes. See the turtle example where the turtle shell was made 100x stiffer than the rest of the turtle.

1.9.5 Can Vega simulate free-flying (unconstrained) deformable objects?

Yes. Simply specify zero constraints when initializing the integrator.

1.9.6 How can I resolve collisions?

Vega does not include collision detection capabilities. However, you can use any external collision detection library, and set the resulting contact forces as external forces in Vega.

1.9.7 I want to embed a triangle mesh into a volumetric mesh, and render the triangle mesh. Does Vega support this?

Yes. It should be noted that all simulations in Vega run on volumetric meshes (except cloth simulations). However, you can transfer (in real-time, or offline) the volumetric mesh deformations to the embedded triangle mesh using the `interpolate` routine in the `VolumetricMesh` class. The technique uses barycentric interpolation. You can see this in the turtle example (and other examples). See the source code of the `interactiveDeformableSimulator` driver.

1.9.8 How can I compute the mass matrix?

You can use `utilities/volumetricMeshUtilities/generateMassMatrix`. It calls a routine from `generateMassMatrix.h` in the `volumetricMesh` library, which you can call directly from your code.

1.9.9 How can I compute the stiffness matrix?

Use `GetTangentStiffnessMatrix` (or `ComputeStiffnessMatrix`) in any of the several provided material classes.

1.9.10 Is there a MS Visual Studio project (solution) file available?

Yes. It has been included since Vega FEM 3.0. Please see the "Microsoft Visual Studio" paragraph in Section 1.1.

1.9.11 What are the physical units used in Vega?

```
Input 3D meshes: meters (m)
Time: seconds (s)
Forces: Newton (N)
Young's modulus: Pa=N/m^2
Poisson's ratio: no unit (dimensionless)
```

1.9.12 How are the invariants I, II, III defined (for isotropic materials)?

Vega follows the definition in the reference [BW08]:

$$I = \mathrm{tr}(C) = \lambda_1^2 + \lambda_2^2 + \lambda_3^2, \tag{2}$$

$$II = \mathrm{tr}(C^2) = \lambda_1^4 + \lambda_2^4 + \lambda_3^4, \tag{3}$$

$$III = \det(C) = \lambda_1^2 \lambda_2^2 \lambda_3^2. \tag{4}$$

Some other references, however, define II as

$$II' = \lambda_1^2 \lambda_2^2 + \lambda_2^2 \lambda_3^2 + \lambda_3^2 \lambda_1^2. \tag{5}$$

It is possible to easily convert from II to II'; the two are related via I^2.

1.9.13 Where can I learn more about FEM, deformable objects, and the methods implemented in Vega?

Our research paper on Vega (published in the Computer Graphics Forum Journal [SSB12]) explains the implemented methods and analyzes the performance of Vega. You can also read the cited papers, and visit the webpage of the SIGGRAPH 2012 course on FEM for deformable object simulation: `http://www.femdefo.org`.

1.9.14 I like Vega. I want to donate funding.

If you want to donate funding, this is of course welcome. You can contact Jernej Barbic at jnb@usc.edu. The funds will be used for further academic research on Vega and deformable object simulation. In any case, we will appreciate if you let us know that you used Vega in your application.

2 Libraries

Below is a listing of selected libraries in the `libraries` folder. The purpose of each library as a whole is described, and more specific information is given on *selected* constituent classes and member functions to highlight important functionality. For subclasses, virtual functions are only re-listed if the subclass implements a previously-abstract function or substantially alters its functionality. Note that a few functions allocate memory which must be deleted by the caller: you can recognize such functions by the ** (pointer to pointer) calling convention. As standard in C/C++, the memory is allocated with malloc/free for the built-in datatypes, and new/delete for all the other datatypes.

2.1 basicAlgorithms

Provides assorted low-level algorithms such as sorting and removing elements. For example, removing multiple elements with specified indices from a vector.

2.2 camera

class SphericalCamera Provides utilities for setting the OpenGL camera based upon a spherical-coordinate system centered at a specified focus point.

```
void Look()
```
Applies the current camera transformation to the active OpenGL matrix. Note that this does not run `glMatrixMode` or `glLoadIdentity`.

```
void MoveRight(double amount)
void MoveUp(double amount)
void MoveIn(double amount)
void ZoomIn(double amount)
```
Move the camera by a user-specified offset. Supply negative values to move/zoom in the opposite direction.

```
void SavePosition(const char * filename)
void LoadPosition(const char * filename)
```
Save or load the camera position to a file. `LoadPosition` prints a warning if the specified file does not exist.

2.3 clothBW

Implements the well-known Baraff-Witkin cloth simulator [BW98]. Input is a triangle mesh with material properties such as stretch, shear and bend resistance. The library can compute both the internal elastic forces and their gradient (tangent stiffness matrix). It computes stretch, shear, and bend forces. For bend forces, the user can choose to use the input angle as the rest bend angle, or use the zero angle. Baraff-Witkin damping is not implemented, but you can use the damping provided by the Vega `integrator` class. In order to timestep the cloth, use the `integratorSparse` library. In this way, you can create cloth animations under any specified external forces. For collisions, you need to use some external library to compute the contact forces on the cloth, and then set them as external forces to the simulator.

class ClothBW   Implements the Baraff-Witkin cloth simulator [BW98], as described above.

```
ClothBW( int numVertices, const double * restPositions, const double * masses,
    int numTriangles, const int * triangles, const int * triangleGroups,
    int numMaterialGroups, const MaterialGroup * materialGroups)
```
Creates the cloth elastic model, from a given triangle mesh. Variable `numVertices` specifies the number of vertices in the triangle mesh, `masses` is an array of length `numVertices` that gives the masses of each vertex, `restPositions` is an array of length $3\times$ `numVertices` and gives the rest positions of the vertices, three values (x, y, z) per each vertex, `triangles` is an integer array of length $3\times$ `numTriangles` giving integer indices of the three vertices forming a triangle, `triangleGroups` is an integer array of length `numTriangles` giving the integer index of the material group to which each triangle belongs, `materialGroups` are arrays that give the scalar stiffness and damping parameters for each material group. All indices in the class are 0-indexed. This constructor does not require triangleUVs input; it computes suitable UVs automatically. The UVs are continuous only within each triangle; the UV map is not global. This is sufficient for isotropic simulation; but cannot accommodate anisotropic effects.

```
ClothBW(int numVertices, const double * restPositions, const double * masses,
    int numTriangles, const int * triangles,
    const double * triangleUVs, const int * triangleGroups,
    int numMaterialGroups, const MaterialGroup * materialGroups)
```
A variant of the constructor where the UVs are not computed automatically, but are provided by the caller. Parameter `triangleUVs` is a double array of length $3 \times 2\times$ `numTriangles`, indicating the *uv* coordinates for every vertex.

```
void ComputeForce(const double * u, double * f, bool addForce=false)
```
Computes the total cloth force on every vertex. Parameter `u` (length $3\times$`numVertices`) gives the displacements of all vertices away from the rest configuration.

```
void ComputeStiffnessMatrix(const double * u, SparseMatrix * K, bool addMatrix=false)
```
Computes the tangent stiffness matrix (gradient of cloth forces). Parameter `u` (length $3\times$`numVertices`) gives the displacements of all vertices away from the rest configuration.

```
void SetComputationMode(const bool mode[4])
```
Specifies what cloth forces and stiffness matrices are to be computed by `ComputeForce` and `ComputeStiffnessMatrix`, `AddForce` and `AddStiffnessMatrix`. This allows users to toggle on/off computation of the stretch/shear forces, bend forces, stretch/shear bend stiffness matrices, and bend stiffness matrix, as follows:

```
mode[0]: computeStretchAndShearForce: yes/no
mode[1]: computeBendForce: yes/no
mode[2]: computeStretchAndShearStiffnessMatrices: yes/no
mode[3]: computeBendStiffnessMatrices: yes/no
```

Default value is **true** for all fields of `mode`.

class ClothBWFromObjMesh   A convenience class that constructs a cloth model from an obj mesh. The mesh need not be triangulated; if it is non-triangular, faces will be split into triangles automatically. There are two member functions `GenerateClothBW`. One sets constant material properties for all the triangles, and the other allows the specification of specific material properties for each obj mesh material (.mtl file).

2.4   configFile

class ConfigFile   Parses values from a text configuration file. The syntax of the configuration file is user-defined. Options can be read as `int`, `bool`, `float`, `double`, `Vec3d`, `char *` (C-string), `std::string` and

`std::vector<std::string>` types. See any of the `.config` files in the `examples` folder for an example of the config file syntax.

`int addOption(const char * optionName, T * destLocation)`

Defined for `T` as `int`, `bool`, `float`, `double`, `char`, `std::string`, `std::vector<std::string>`. Adds a mandatory option with name `optionName`. When the config file is later parsed, any value found for this option is written to the variable pointed to by `destLocation`. If no value for this option is found, the parse is considered to have failed. Returns a non-zero value if the option has already been defined.

`int addOptionOptional(const char * optionName, T * destLocation, T defaultvalue)`
`int addOptionOptional(const char * optionName,`
`    char * destLocation, const char * defaultvalue)`

Adds an optional option with name `optionName`. When the config file is later parsed, the variable pointed to by `destLocation` is set to the value found in the file, or to `defaultValue` if no value is set. Returns a non-zero value if the option has already been defined.

`int parseOptions(const char * filename)`

Parses the options in file `filename`, writing the option values it reads to the appropriate variables. A non-zero value is returned if any mandatory options are not specified.

2.5 constrainedDOFs

This library was named `insertRows` before version 4.0. Provides utilities for the insertion and removal of elements from dense 1D arrays. This is useful, for example, when constraining (fixing) vertices in a deformable simulation.

`void InsertDOFs(int mFull, double * xConstrained, double * x,`
`    int numFixedDOFs, int * fixedDOFs, int oneIndexed=0)`

Copies `xConstrained` into `x`, and then inserts zeros into `x`. Zeroes are placed among the elements at each index indicated in `fixedDOFs`, until the total desired size `mFull` of the output array is reached. `x` is assumed to be already allocated.

`void RemoveDOFs(int mFull, double * xConstrained, double * x,`
`    int numFixedDOFs, int * fixedDOFs, int oneIndexed=0)`

Copies `x` into `xConstrained`, and then removes the elements at the indices given in `fixedDOFs` from `xConstrained`. `xConstrained` is assumed to be already allocated.

`void FullDOFsToConstrainedDOFs(int mFull, int numDOFs,`
`    int * DOFsConstrained, int * DOFs, int numFixedDOFs,`
`    int * fixedDOFs, int oneIndexed=0)`

Translates indices of elements in an unreduced array to the indices which would contain those same elements in the reduced array produced by calling `RemoveDOFs` with `fixedDOFs`. Input is `DOFs` and output is `DOFsConstrained`. Writes to `DOFsConstrained[i]` the index in the reduced array that would give the element at index `DOFs[i]` in the unreduced array.

2.6 corotationalLinearFEM

class CorotationalLinearFEM Implements the corotational linear finite element model described in [MG04]. The class can also compute the exact tangent stiffness matrix for tet meshes; the implementation is described in [Bar12].

`CorotationalLinearFEM(VolumetricMesh * mesh)`
Initializes the model based upon an input volumetric mesh.

```
void ComputeElementEnergyAndForceAndStiffnessMatrix(int elementID,
    const double * vertexDisplacements, double * energy,
    double * elementInternalForces, double * elementStiffnessMatrix, int warp=1)
```
Computes the energy, internal forces and stiffness matrix on one volumetric mesh element given a vector `vertexDisplacements` of the displacements for the vertices of the volumetric mesh. If `energy`, `internalForces` or `stiffnessMatrix` is NULL, the function does not calculate or return the corresponding information. The `warp` parameter controls whether the simulation "warps" stiffnesses and therefore supports large deformations. When warping is enabled (`warp=1` or `warp=2`), the implementation supports large deformations. The default is `warp=1`, in which case the simulation uses the approximate tangent stiffness matrix as described in [MG04]. For `warp=2`, the class computes the exact tangent stiffness matrix; our implementation is described in [Bar12]. Such an exact matrix has better simulation properties; however, it only supports tet meshes and requires approximately 1.5x the computation time of the approximate matrix. If warping is disabled (`warp=0`), one obtains the standard linear FEM simulation [Sha90]. Such simulation runs faster than the warped simulations because it timesteps the linear equation $M\ddot{u} + D\dot{u} + Ku = f$. It is only accurate under small displacements.

2.7   distanceField

Calculates a signed or unsigned distance field for 3D geometry. The distance field can be full (computed on a regular 3D lattice), or only computed in a narrow band around the geometry.

class DistanceField : public DistanceFieldBase   Stores and computes a full distance field. If Intel TBB is enabled, it computes in parallel.

```
virtual int computeUnsignedField(ObjMesh * objMesh, int resolutionX, int resolutionY,
    int resolutionZ, int maxTriCount=15, int maxDepth=10, int zMin = -1, int zMax = -1)
virtual int computeSignedField(ObjMesh * objMesh, int resolutionX, int resolutionY,
    int resolutionZ, int maxTriCount=15, int maxDepth=10, int zMin = -1, int zMax = -1)
```
Computes a full unsigned or signed distance field. The field has the following number of grid points: (`resolutionX`+1) $\times$ (`resolutionY`+1) $\times$ (`resolutionZ`+1). The input obj mesh is assumed to be manifold and self-intersection-free when computing a signed distance field.

```
virtual inline float distance(int i, int j, int k)
```
Returns distance field value at grid vertex (i, j, k).
Each of i, j, k must be an integer from $[0, $ `resolution{X,Y,Z}`$]$.

```
virtual float  distance(Vec3d pos, int constrainToBox=0)
```
Computes distance at an arbitrary position.

class DistanceFieldNarrowBand : public DistanceFieldBase   Stores and computes a narrow band distance field. Not all distances at grid vertices are computed. Only the grid vertices that are within the prescribed distance to the geometry are calculated and stored, so this representation is much more efficient that a full distance field.

```
virtual int computeUnsignedField(ObjMesh * objMesh, int resolutionX, int resolutionY,
    int resolutionZ, double bandWidth, int maxTriCount=15, int maxDepth=10)
virtual int computeSignedField(ObjMesh * objMesh, int resolutionX, int resolutionY,
    int resolutionZ, double bandWidth, int maxTriCount=15, int maxDepth=10)
```
Computes an unsigned and signed distance field in a narrow band. The parameter `bandWidth` is given in absolute units. Only distances at grid vertices that are smaller than `bandWidth` are computed and stored. Just like `DistanceField`, this class also assumes that the mesh is manifold and self-intersection-free when computing a signed distance field.

class SignedDistanceFieldFromPolygonSoup Calculates a reasonable signed distance field from "polygon soup" geometry, as described in [XB14]. "Polygon soup" geometry here refers to arbitrary, unorganized 3D triangular geometry that may self-intersect, be duplicated, or have cracks or holes. For usage, see the comments in `signedDistanceFieldFromPolygonSoup.h`.

class DistanceFieldCreator Provides an easy interface to various distance field generation methods.

```
DistanceField * ComputeDistanceField(int resolutionX, int resolutionY, int resolutionZ,
    int calculateSignedField, SignedFieldCreationMode mode, double sigma,
    int subtractSigma = 1, bool computeVoronoiDiagram = false,
    int maxTriCount = 15, int maxDepth = 10, int closestPointFlag = 0,
    const char * precomputedUnsignedFieldFilename = NULL)
```
Computes a full distance field. An unsigned field is computed if `calculateSignedField` is 0, and a signed field otherwise. For signed field calculation, the `mode` variable specifies which signed field generation method to use. If Intel TBB is enabled, it computes in parallel. We provide three methods:

- BASIC: assumes that the input mesh is manifold and self-intersection-free. It uses the methods in `DistanceField` class to compute a signed distance field.

- POLYGONSOUP: uses the "signed distance field from polygon soup" pipeline [XB14] to calculate a meaningful signed distance field even for non-manifold and self-intersected meshes.

- AUTO: uses BASIC if mesh is manifold, otherwise POLYGONSOUP.

```
DistanceFieldNarrowBand * ComputeDistanceFieldNarrowBand(int resolutionX, int resolutionY,
    int resolutionZ, double bandWidth, int calculateSignedField,
    SignedFieldCreationMode mode, double sigma, int subtractSigma = 1,
    int maxTriCount = 15, int maxDepth = 10,
    const char * precomputedUnsignedFieldFilename = NULL)
```
Computes a narrow band distance field. Its parameters are the same as `ComputeDistanceField`. In BASIC mode, the `DistanceFieldNarrowBand` class will be used for the calculation.

2.8 elasticForceModel

Note: deprecated in version 4.0. New designs should use the `StencilForceModel`. Provides implementation of the `ForceModel` base class for the deformable models supported by Vega. This makes it possible to use these deformable models with the integrators in Vega (`integrator` library).

Starting from version 4.0, Vega uses `stencilForceModel` to replace the role of `elasticForceModel`. The class `elasticForceModel` can now only be used for single-core computation of elastic forces and stiffnesses matrices. For multicore, use `StencilForceModel`. This new structure greatly improves multi-threading efficiency and management, and also eliminates the need to implement a "MT" class for each force model. Therefore, the "MT" classes (e.g., `corotationalLinearFEMMT`) have been removed in version 4.0.

2.9 exactArithmetic

A wrapper for CGAL's exact-predicates-exact-constructions kernel. The class `ER` in `exactArithmetic` inherits this CGAL kernel. You can replace this wrapper with other exact arithmetic libraries.

2.10 forceModel

This library provides the abstract base class for a force model used in the `integratorSparse` library, i.e., a "black-box" function $u \mapsto f_{\text{int}}(u)$ and its gradient in

$$M\ddot{u} + \big(\alpha M + \beta K(u) + D\big)\dot{u} + f_{\text{int}}(u) = f_{\text{ext}}.$$

Any dynamical system described by such a differential equation can then be timestepped by the `integrator` library, by providing an implementation of f_{int} and its gradient, in a class derived from `ForceModel`.

For deformable simulations, class `ForceModel` connects integrators to internal forces and tangent stiffness matrix calculator classes. This allows the different deformable models in Vega to expose their internal forces and stiffness matrices to the integrator in a uniform way. Derived classes that implement f_{int} and its gradient for the deformable models in Vega can be found in the library `stencilForceModel`. The reason for why library `forceModel` is separated from `stencilForceModel` is so that the `integratorSparse` library can be compiled and used independently from any deformable materials in Vega. Similarly, one can compile and use `forceModel` and `stencilForceModel` independently of `integratorSparse`, making it possible to timestep the Vega deformable models with externally provided integrators.

class ForceModel Abstract base class for a force model $f_{\mathrm{int}}(u)$ whose gradient (typically the tangent stiffness matrix) is a *sparse* matrix. All deformable models in Vega fall into this category. Note: dense gradient matrices occur in applications involving model reduction.

    int Getr()
Returns `r`, the number of object degrees of freedom (typically three times the number of vertices or particles). Note that the `r` member variable must be set by a subclass.

    virtual void GetInternalForce(double * u, double * internalForces) = 0
The internal forces arising from vertex displacements `u` are written to the array `internalForces`. Must be implemented by derived classes.

    virtual void GetTangentStiffnessMatrixTopology(SparseMatrix **
       tangentStiffnessMatrix) = 0
Allocates (`new`) a `SparseMatrix` with the correct pattern of non-zero entries to hold a stiffness matrix for this material, and write the matrix pointer to `*tangentStiffnessMatrix`. Must be implemented by derived classes.

    virtual void GetTangentStiffnessMatrix(double * u,
       SparseMatrix * tangentStiffnessMatrix) = 0
The tangent stiffness matrix arising from vertex displacements `u` is written to the previously allocated `tangentStiffnessMatrix`. Must be implemented by derived classes.

    virtual void GetForceAndMatrix(double * u, double *
       internalForces, SparseMatrix * tangentStiffnessMatrix)
Writes out the internal forces and stiffness matrix arising from displacements `u`, using the implementations of the functions above. May be overloaded by derived classes.

2.11   getopts

    int getopts(int argc, char **argv, opt_t opttable[])
Parses the `argc` and `argv` from a main function and extracts any specified option values. Modified from public domain code by Paul Edwards.

class CommandLineParser   A C++ implementation of `getopts` .

2.12   glslPhong

Implements per-pixel (Phong) lighting using a GLSL shader.

2.13   graph

class Graph   Stores an undirected graph. The vertices of the graph are represented by the integers from 0 to the number of vertices minus one, and the graph structure is constant after being set in the constructor.

`Graph(int numVertices, int numEdges, const int * edges, int sortEdgeVertices=1)`

Initializes the graph, giving it `numVertices` vertices and `numEdges` edges. The input array of integers `edges` encodes the edges as subsequent pairs of vertex indices.

`int IsNeighbor(int vtx1, int vtx2)`

Returns 0 if vertices `vtx1` and `vtx2` are not neighbors, and returns 1 plus the index of vtx2 in the list of `vtx1`'s neighbors otherwise.

2.14   hashTable

Implements a simple 1D hash table. Keys are of type "unsigned int", and datatype can be arbitrary (templated).

2.15   imageIO

Loads and saves PNG, TIFF and JPEG, TGA, PPM file formats. In order to use PNG, JPEG or TIFF, you must enable them in the file "imageFormats.h", and recompile the code, and then link against libpng, libjpeg and libtiff libraries, respectively. PPM and TGA are built-in. They need not be enabled (they are always enabled) and require no linking against external libraries.

class ImageIO   Performs the functionality described above.

2.16   immersionMesher

class ImmersionMesher   Given an input tet mesh that covers the space occupied by an input manifold self-intersecting triangle mesh, computes a tet mesh with properly duplicated tets to embed the triangle mesh. The triangle mesh is assumed to be manifold and without degenerate triangles. The method duplicates the tets and tet vertices to follow the topology of the triangle mesh. The implementation follows the following paper:

```
Yijing Li, Jernej Barbic:
Immersion of Self-Intersecting Solids and Surfaces,
ACM Transactions on Graphics 37(4) (SIGGRAPH 2018)
```

```
void run( const TriMeshGeo & triMesh, const TetMeshGeo & tetMesh,
   std::vector<TetMeshGeo> & outputTetMeshes,
   std::vector<BarycentricCoordinates> & embeddingWeights,
   std::vector<std::vector<TriMeshGeo>> * cellMeshes = nullptr,
   std::vector<std::vector<BarycentricCoordinates>> * cellMeshWeights = nullptr )
```

The main routine. It computes an "unglued" tet mesh by properly duplicating the tets of the input tet mesh, and embeds the input triangle mesh into it (please see paper). Input `triMesh` and `tetMesh` are the triangle and tet mesh geometry. `triMesh` is manifold, non-degenerate and may be self-intersecting. `TetMesh` should cover the space occupied by the triangle mesh. Each output of the function is an std::vector because there may be multiple tet mesh outputs, due to ambiguities in the immersion (please see paper). `outputTetMeshes` stores output tet meshes, and `embeddingWeights` stores barycentric coordinates of `triMesh` embedded into each output tet mesh. Optionally for visualization purposes, `cellMeshes` gives the triangle mesh of each cell and `cellMeshWeights` gives the barycentric coordinates to embed `cellMeshes` into the output tet meshes.

2.17   integrator

This is the "base" library for numerical integration in Vega. Together with "derived" libraries `integratorSparse` and `integratorDense`, it provides several implicit and explicit integrators to solve equations of the form

$$M\ddot{u} + \big(\alpha M + \beta K(u) + D\big)\dot{u} + f_{\text{int}}(u) = f_{\text{ext}}. \tag{6}$$

For a broader discussion of this equation, please see [SSB12]. The integrator library itself only consists of abstract base classes. You must use either `integratorSparse` or `integratorDense` for the actual integration. The sparse library (`integratorSparse`) is designed for systems (6) where M, D, K are (large) sparse matrices, e.g., geometrically complex deformable objects (without model reduction). The dense library (`integratorDense`), in turn, is designed for systems (6) where M, D, K are dense matrices, e.g., for model reduction simulations.

class IntegratorBase   Serves as a base class to all the integrators. The class stores displacements, velocities, and accelerations for the simulation internally, and provides access to modify these buffers, but offers no actual timestepping capabilities; this is deferred to derived classes. The class also holds Rayleigh damping coefficients that specify how much the mass and stiffness matrix contribute to the damping matrix.

    IntegratorBase(int r, double timestep,
        double dampingMassCoef=0.0, double dampingStiffnessCoef=0.0)
Sets the number of degrees of freedom `r`, timestep `timestep`, and coefficients for how much to add the mass matrix and stiffness matrix to the damping matrix. Allocates (`malloc`) a number of internal buffers of `r doubles` for holding the current displacement, velocity, internal forces, and so on.

    virtual ~IntegratorBase()
Frees (`free`) the internal buffers.

    virtual void ResetToRest()
Sets the internal position, velocity, and acceleration buffers to zero.

    virtual int SetState(double * q, double * qvel=NULL) = 0
Resets internal position buffer to the values in `q`, and does likewise for internal velocity buffer if `qvel` is not NULL. Returns 0 if successful, 1 otherwise.

    void SetqState(const double * q, const double * qvel=NULL, const double * qaccel=NULL)
    void GetqState(double * q, double * qvel=NULL, double * qaccel=NULL)
Copies external position, velocity, and/or acceleration buffers to the internal buffers, or vice versa. Each buffer is copied only if the pointer to the external buffer is not NULL.

    void SetExternalForces(double * externalForces)
    void AddExternalForces(double * externalForces)
    void GetExternalForces(double * externalForces)
Sets or adds the values from `externalForces` to the external forces buffer, or writes the external forces buffer to `externalForces`.

    virtual void SetTimestep(double timestep)
    double GetTimestep()
Sets or returns the timestep value.

    virtual int DoTimestep() = 0
Performs a timestep of the simulation, given the current values in the internal/external forces, position, velocity, and acceleration buffers. The resulting position, velocity, and acceleration are saved to these buffers. Returns 0 if and only if the timestep is completed without error.

23

2.18 integratorSparse

This library can timestep systems (6) where M, D, K are (large) sparse matrices, e.g., geometrically complex deformable objects (without model reduction). This is the "core" integrator capability in Vega. For model reduction, see `integratorDense`.

class IntegratorBaseSparse : public IntegratorBase A base class for integrators for dynamical systems characterized by (large) sparse matrices, such as geometrically complex 3D elasticity. This is the main integrator type in Vega. These integrators use a `ForceModel` class to obtain the internal forces and stiffness matrices. Stiffness matrices are sparse and stored using the `SparseMatrix` class. `IntegratorBaseSparse` stores two sparse matrices: the mass matrix and the damping matrix to be applied in addition to mass- and stiffness-based damping specified in `IntegratorBase`.

```
IntegratorBaseSparse(int r, double timestep, SparseMatrix * massMatrix,
    ForceModel * forceModel, int numConstrainedDOFs=0, int * constrainedDOFs=NULL,
    double dampingMassCoef=0.0, double dampingStiffnessCoef=0.0)
```
Initializes the number of degrees of freedom `r`, the timestep, the mass matrix, which degrees of freedom are constrained, and the `ForceModel`. The internal damping `SparseMatrix` is set to zero.

```
virtual void SetForceModel(ForceModel * forceModel)
```
Sets the `ForceModel` object to `forceModel`. This makes it possible to change the force model at runtime.

```
virtual void SetDampingMatrix(SparseMatrix * dampingMatrix)
```
Sets the damping matrix D to `dampingMatrix`. This matrix will be added to the matrix produced according to the mass and stiffness damping coefficients to get the total damping matrix for the simulation.

```
virtual double GetForceAssemblyTime()
virtual double GetSystemSolveTime()
```
Return the time taken in `DoTimestep` to generate the force/stiffness values and to solve the linear system while timestepping, respectively. It is up to subclasses to calculate these values in `DoTimestep`.

```
virtual double GetKineticEnergy()
virtual double GetTotalMass()
```
Calculates and returns the kinetic energy or total mass of the simulation, based upon the mass matrix and the velocity buffer of the class.

class ImplicitNewmarkSparse : public IntegratorBaseSparse Implements implicit Newmark integration, and expands `IntegratorBaseSparse` with features common to Newmark-style integrators.

```
ImplicitNewmarkSparse(int r, double timestep,
    SparseMatrix * massMatrix, ForceModel * forceModel,
    int numConstrainedDOFs=0, int * constrainedDOFs=NULL,
    double dampingMassCoef=0.0, double dampingStiffnessCoef=0.0,
    int maxIterations = 1, double epsilon = 1E-6,
    double NewmarkBeta=0.25, double NewmarkGamma=0.5,
    int numSolverThreads=0)
```
Initializes `maxIterations`, `epsilon`, `NewmarkBeta`, `NewmarkGamma`, and `numSolverThreads` parameters, and forwards the other parameters to the `IntegratorBaseSparse` constructor. For PARDISO and SPOOLES solvers, we found best performance with 2-3 threads; more threads usually decreased performance.

```
virtual void SetTimestep(double timestep)
```
Sets the timestep.

```
virtual int SetState(double * q, double * qvel=NULL)
```
Sets the position (and optionally the velocity) buffer, and calculates the acceleration buffer accordingly using implicit Newmark, assuming no external force. Returns 0 if successful, 1 otherwise.

```
virtual int DoTimestep()
```
Runs a timestep of implicit Newmark, updating the internal position, velocity, and acceleration buffers accordingly. Returns 0 if and only if the timestep is completed successfully.

```
void SetNewmarkBeta(double NewmarkBeta)
void SetNewmarkGamma(double NewmarkGamma)
```
Set the value of the beta and gamma parameters, respectively, for Newmark integrators. No checking is performed to see if these values are in the appropriate range.

```
void UseStaticSolver(bool useStaticSolver)
```
Sets whether to use a static solver. The class defaults to dynamic.

class ImplicitBackwardEulerSparse : public ImplicitNewmarkSparse   Implements implicit backward Euler integration [BW98].

```
ImplicitBackwardEulerSparse(int r, double timestep,
   SparseMatrix * massMatrix,
   ForceModel * forceModel,
   int numConstrainedDOFs=0, int * constrainedDOFs=NULL,
   double dampingMassCoef=0.0, double dampingStiffnessCoef=0.0,
   int maxIterations = 1, double epsilon = 1E-6,
   int numSolverThreads=0)
```
Initializes the integrator.

```
virtual int SetState(double * q, double * qvel=NULL)
```
Sets the position (and optionally the velocity) buffer based on the parameters, and calculates the appropriate acceleration buffer values using implicit Euler, assuming no external forces. Returns 0 if successful, 1 otherwise.

```
virtual int DoTimestep()
```
Runs a timestep of implicit Euler, and updates the internal position, velocity, and acceleration buffers accordingly. Returns 0 if successful, 1 otherwise.

class CentralDifferencesSparse : public IntegratorBaseSparse   Implements the explicit central differences integrator.

```
CentralDifferencesSparse(int numDOFs, double timestep,
   SparseMatrix * massMatrix,
   ForceModel * forceModel,
   int numConstrainedDOFs=0, int * constrainedDOFs=NULL,
   double dampingMassCoef=0.0, double dampingStiffnessCoef=0.0,
   int tangentialDampingMode=1, int numSolverThreads=0)
```
Initializes the integrator settings via the `IntegratorBaseSparse` constructor, and selects the number of threads to use for the sparse linear solver. Tangential damping mode controls how often the Rayleigh damping matrix is recomputed. This damping matrix depends on the tangent stiffness matrix, which changes in time. When 0, the damping matrix is never updated. The system matrix is then constant, so one can factor it only once. However, this is not recommended for large deformations as it leads to damping artifacts. When tangentialDampingMode > 0, the damping matrix is updated every "tangentialDampingMode"th timestep. Default is 1, i.e., update at every timestep.

```
virtual int SetState(double * q, double * qvel=NULL)
```
Sets the position (and optionally velocity) buffers. Always returns 0.

```
virtual int DoTimestep()
```
Runs a timestep of explicit central differences, and updates the internal position, velocity and acceleration buffers accordingly. Returns 0 if successful, 1 otherwise.

class EulerSparse : public IntegratorBaseSparse Implements explicit Euler integration, with a flag to enable symplectic Euler integration. Because this class never forms the tangent stiffness matrix, damping is controlled via mass-based damping only.

```
EulerSparse(int r, double timestep, SparseMatrix * massMatrix,
    ForceModel * forceModel,
    int symplectic=0, int numConstrainedDOFs=0,
    int * constrainedDOFs=NULL, double dampingMassCoef=0.0)
```
Initializes the integrator settings via the `IntegratorBaseSparse` constructor, and sets whether to perform symplectic integration.

```
virtual int SetState(double * q, double * qvel=NULL)
```
Sets the position (and optionally velocity) buffers based on the parameters, and sets the acceleration buffer via explicit (or symplectic) Euler assuming no external forces. Returns 0 if successful, 1 otherwise.

```
virtual int DoTimestep()
```
Runs a timestep of explicit (or symplectic) Euler, and updates the internal position, velocity, and acceleration buffers accordingly. Returns 0 if successful, 1 otherwise.

2.19   integratorDense

This library can timestep systems (6) where M, D, K are dense matrices. While this functionality is general-purpose, its main purpose in Vega is model reduction. This class largely parallels `integratorSparse`, with the class names simply renamed from "Sparse" to "Dense". For standard (non-model reduction) simulation in Vega, see `integratorSparse`.

class IntegratorBaseDense : public IntegratorBase A base class for integrators for dynamical systems characterized by dense matrices. These integrators use a `ReducedForceModel` class to obtain the internal forces and stiffness matrices. Stiffness matrices are dense and stored in column-major format.

```
IntegratorBaseDense(int r, double timestep, double * massMatrix,
    ReducedForceModel * reducedForceModel,
    double dampingMassCoef=0.0, double dampingStiffnessCoef=0.0)
```
Initializes the number of degrees of reduced freedom `r`, the timestep, the mass matrix, and the `ReducedForceModel`. The internal damping is set to zero.

```
void SetReducedForceModel(ReducedForceModel * reducedForceModel)
```
Sets the `ReducedForceModel` object to `reducedForceModel`. This makes it possible to change the force model at runtime.

```
virtual void SetMassMatrix(double * massMatrix)
```
Sets the reduced mass matrix to `massMatrix`. This makes it possible to change the mass matrix at runtime.

```
virtual double GetForceAssemblyTime()
virtual double GetSystemSolveTime()
```

Return the time taken in `DoTimestep` to generate the force/stiffness values and to solve the linear system while timestepping, respectively. It is up to subclasses to calculate these values in `DoTimestep`.

```
virtual double GetKineticEnergy()
virtual double GetTotalMass()
```
Calculates and returns the kinetic energy or total mass of the simulation, based upon the mass matrix and the velocity buffer of the class.

class ImplicitNewmarkDense : public IntegratorBaseDense Implements implicit Newmark integration, and expands `IntegratorBaseDense` with features common to Newmark-style integrators. This is the integrator used in [BJ05].

```
ImplicitNewmarkDense(int r, double timestep,
    double * massMatrix, ReducedForceModel * reducedForceModel,
    solverType solver=positiveDefiniteMatrixSolver,
    double dampingMassCoef=0.0, double dampingStiffnessCoef=0.0,
    int maxIterations = 1, double epsilon = 1E-6, double NewmarkBeta=0.25,
    double NewmarkGamma=0.5)
```
Initializes `maxIterations`, `epsilon`, `NewmarkBeta`, and `NewmarkGamma` parameters, and forwards the other parameters to the `IntegratorBaseDense` constructor. There are three choices for the dense solver: positive-definite, symmetric, general. Positive-definite solver is the most common choice (and cca 2x faster than other options). However, if the system matrix becomes singular (rare; only extreme deformations), an error will be issued. To specify the solver, use one of the three choices: `solverType::generalMatrixSolver`, `solverType::symmetricMatrixSolver`, `solverType::positiveDefiniteMatrixSolver`.

```
virtual void SetTimestep(double timestep)
```
Sets the timestep.

```
virtual int SetState(double * q, double * qvel=NULL)
```
Sets the position (and optionally the velocity) buffer, and calculates the acceleration buffer accordingly using implicit Newmark, assuming no external force. Returns 0 if successful, 1 otherwise.

```
virtual int DoTimestep()
```
Runs a timestep of implicit Newmark, updating the internal position, velocity, and acceleration buffers accordingly. Returns 0 if and only if the timestep is completed successfully.

```
void SetNewmarkBeta(double NewmarkBeta)
void SetNewmarkGamma(double NewmarkGamma)
```
Set the value of the beta and gamma parameters, respectively, for Newmark integrators. No checking is performed to see if these values are in the appropriate range.

```
void UseStaticSolver(bool useStaticSolver)
```
Sets whether to use a static solver. The class defaults to dynamic.

class ImplicitBackwardEulerDense : public IntegratorBaseDense Implements implicit backward Euler integration [BW98], for dense stiffness matrices.

```
ImplicitBackwardEulerDense(int r, double timestep,
    double * massMatrix,
    ReducedForceModel * reducedForceModel,
    solverType solver=positiveDefiniteMatrixSolver, double dampingMassCoef=0.0,
    double dampingStiffnessCoef=0.0, int maxIterations = 1,
    double epsilon = 1E-6)
```
Sets the integrator parameters with the `ImplicitNewmarkDense` constructor.

```
virtual int SetState(double * q, double * qvel=NULL)
```
Sets the position (and optionally the velocity) buffer based on the parameters, and calculates the appropriate acceleration buffer values using implicit Euler, assuming no external forces. Returns 0 if successful, 1 otherwise.

```
virtual int DoTimestep()
```
Runs a timestep of implicit Euler, and updates the internal position, velocity, and acceleration buffers accordingly. Returns 0 if successful, 1 otherwise.

class CentralDifferencesDense : public IntegratorBaseDense   Implements the explicit central differences integrator.

```
CentralDifferencesDense(int numDOFs, double timestep,
    double * massMatrix,
    ReducedForceModel * reducedForceModel,
    double dampingMassCoef=0.0, double dampingStiffnessCoef=0.0,
    int tangentialDampingMode=1)
```
Initializes the integrator settings via the `IntegratorBaseDense` constructor. Tangential damping mode controls how often the Rayleigh damping matrix is recomputed. This damping matrix depends on the tangent stiffness matrix, which changes in time. When 0, the damping matrix is never updated. The system matrix is then constant, so one can factor it only once. However, this is not recommended for large deformations as it leads to damping artifacts. When tangentialDampingMode > 0, the damping matrix is updated every "tangentialDampingMode"th timestep. Default is 1, i.e., update at every timestep.

```
virtual int SetState(double * q, double * qvel=NULL)
```
Sets the position (and optionally velocity) buffers. Always returns 0.

```
virtual int DoTimestep()
```
Runs a timestep of explicit central differences, and updates the internal position, velocity and acceleration buffers accordingly. Returns 0 if successful, 1 otherwise.

2.20   isotropicHyperelasticFEM

Provides classes to calculate the energy, internal forces and stiffness matrices of a variety of invertible material methods.

class IsotropicHyperelasticFEM   Provides an implementation of the invertible FEM method presented in [ITF04]. Also computes the tangent stiffness matrix as presented in [TSIF05]. Supports isotropic hyperelastic materials with energy defined in terms of the three invariants I, II, III. A few such example materials are provided, and the user can define her own custom materials simply by deriving from `IsotropicMaterial`.

```
IsotropicHyperelasticFEM(TetMesh * tetMesh, IsotropicMaterial * isotropicMaterial,
    double inversionThreshold=-DBL_MAX)
```
Initializes the invertible FEM method. Before creating this class, you must first create the tet mesh, and create an instance of the `IsotropicMaterial` material (e.g., `NeoHookeanIsotropicMaterial`). Parameter `inversionThreshold` controls when the inversion prevention mechanism activates. If the principal stretches are smaller than the `inversionThreshold`, they will be clamped to that, which will generally prevent element inversion. For example, a typical `inversionThreshold` value would be 0.1. By default, inversion handling is disabled (`inversionThreshold`$=-\infty$). Values of 1.0 or higher should not be used. The `isotropicMaterial` must be pre-created: you can use one of our provided classes, or derive from the `IsotropicMaterial` to create your own custom isotropic hyperelastic materials.

The material properties are determined as follows. If the `isotropicMaterial` is of the "Homogeneous" kind (for example, `HomogeneousNeoHookeanIsotropicMaterial`), then the material properties are homogeneous and are specified by `isotropicMaterial`; material properties in the `tetMesh` are ignored (only geometry is used). If, however, the `isotropicMaterial` is not of the "Homogeneous" kind (for example, `NeoHookeanIsotropicMaterial`), then the material properties are such as specified in the `tetMesh` and `isotropicMaterial` class (and may be non-homogeneous).

Internally, the constructor reads the mesh `tetMesh`, pre-computes the area weighted vertex normals, the inverse of the Dm (as in [ITF04]), and the derivative of the deformation gradient with respect to the vertex displacements.

```
void ComputeElementLocalEnergyAndForceAndMatrix(int elementID,
    const double * vertexDisplacements, double * energy,
    double * elementInternalForces, double * elementStiffnessMatrix)
```

Computes the non-linear elastic strain energy, internal forces and stiffness matrix on one volumetric mesh element given a vector `vertexDisplacements` of the displacements for the vertices of the tetrahedral mesh. If `energy`, `internalForces` or `stiffnessMatrix` is NULL, the function does not calculate or return the corresponding information. The energy and forces are calculated based upon the principal stretches of the SVD-diagonalized deformation gradients of the mesh element. The stiffness matrix is calculated based upon the principal stretches and rotation matrices from the SVD diagonalization of the deformation gradient of the element.

class IsotropicMaterial   Serves as an abstract base class for isotropic hyperelastic materials which define energy in terms of the three invariants I, II, III. This classes is then derived into classes that implement concrete isotropic materials. The class represents the material of a single element (tet or cube). The definition of the invariants used in Vega FEM is as follows:

$$I = \operatorname{tr}(C) = \lambda_1^2 + \lambda_2^2 + \lambda_3^2, \tag{7}$$

$$II = \operatorname{tr}(C^2) = \lambda_1^4 + \lambda_2^4 + \lambda_3^4, \tag{8}$$

$$III = \det(C) = \lambda_1^2 \lambda_2^2 \lambda_3^2, \tag{9}$$

where λ_i are the principal stretches (singular values of the deformation gradient F), and $C = F^T F$ is the right Cauchy-Green deformation tensor. Note that in the literature, there are at least two ways to define II. The above invariants are defined in [BW08]. Some other references define the second invariant as

$$II' = \lambda_1^2 \lambda_2^2 + \lambda_2^2 \lambda_3^2 + \lambda_3^2 \lambda_1^2. \tag{10}$$

Such an invariant can be easily converted to the invariants used by Vega, and back, $II' = (I^2 - II)/2$.

```
virtual double ComputeEnergy(int elementIndex, double * invariants) = 0
```
Returns the energy of the element given the three invariants `invariants[0]`, `invariants[1]`, and `invariants[2]`. The integer variable elementIndex makes it possible to make the material properties vary from element to element (heterogeneous material properties).

```
virtual void ComputeEnergyGradient(int elementIndex,
    double * invariants, double * gradient) = 0
```
Given the three invariants `invariants[0]`,`invariants[1]`,`invariants[2]` computes the derivative of the energy with respect to the invariants and writes the result to `gradient[0]`,`gradient[1]`,`gradient[2]`.

```
virtual void ComputeEnergyHessian(int elementIndex,
    double * invariants, double * hessian) = 0
```
Given the three invariants `invariants[0]`,`invariants[1]`,`invariants[2]` computes the second derivative of the energy with respect to the invariants and writes the result to `hessian` where `hessian[0]` = $\partial^2 \Psi / \partial I^2$, `hessian[1]` = $\partial^2 \Psi / \partial I \partial II$, `hessian[2]` = $\partial^2 \Psi / \partial I \partial III$, `hessian[3]` = $\partial^2 \Psi / \partial II^2$, `hessian[4]` = $\partial^2 \Psi / \partial II \partial III$, and `hessian[5]` = $\partial^2 \Psi / \partial III^2$

To define your own isotropic hyperelastic material, simply derive from `IsotropicMaterial` and implement functions `ComputeEnergy`, `ComputeEnergyGradient`, `ComputeEnergyHessian`.

class IsotropicMaterialWithCompressionResistance : public IsotropicMaterial   This abstract class is the same as `IsotropicMaterial`, except it adds the ability for the derived classes to use compression resistance. Compression resistance is an extra energy term which prevents the material from undergoing large compression. Such a capability makes the material more volume preserving and improves simulation stability. All the derived classes in Vega implement compression resistance as explained in [KTY09]. Compression resistance is internally scaled with the material stiffness. Therefore, the value of compression resistance has (approximately) equal effect on the simulation regardless of stiffness. In practice, using compression resistance improves simulation quality and stability. It also produces better looking simulations under large deformations. Its use is highly recommended.

```
IsotropicMaterialWithCompressionResistance(int enableCompressionResistance=0)
```

class StVKIsotropicMaterial : public IsotropicMaterialWithCompressionResistance   StVK material that supports spatially varying material parameters. Provides functions to compute the energy, gradient, and Hessian of the StVK material given the three invariants I, II, III. The implemented StVK material is described in [BW08], page 158.

```
StVKIsotropicMaterial(TetMesh * tetMesh, int enableCompressionResistance=0,
    double compressionResistance=0.0)
```
The material parameters are read from the tet mesh, and are cached in the constructor. Parameter `compressionResistance` must be non-negative. The larger the value, the more the material resists compression. Value of 0 means that there is no resistance; this is equivalent to setting `enableCompressionResistance` to 0. A reasonable value of `compressionResistance` is 500.0. High values will require smaller timesteps, or else the simulation may go unstable. Largest values we typically tried in practice were in 10,000s, e.g., 50,000. The constructor throws an exception if the provided tet mesh does not consist of materials specifying E, ν.

class HomogeneousStVKIsotropicMaterial : public IsotropicMaterialWithCompressionResistance   StVK material with homogeneous (spatially global) material parameters. Provides functions to compute the energy, gradient, and Hessian of the StVK material given the three invariants I, II, III. Compression resistance is as described in [KTY09].

```
HomogeneousStVKIsotropicMaterial(double E, double nu,
    int enableCompressionResistance=0, double compressionResistance=0.0)
```
Initializes the Young's modulus `E` and the Poisson ratio `nu` of the material.

class NeoHookeanIsotropicMaterial : public IsotropicMaterialWithCompressionResistance   Neo-Hookean material that supports spatially varying material parameters. Provides functions to compute the energy, gradient, and Hessian of the Neo-Hookean material given the three invariants I, II, III. The implemented neo-Hookean material is described in [BW08], page 162.

```
NeoHookeanIsotropicMaterial(TetMesh * tetMesh,
    int enableCompressionResistance=0, double compressionResistance=0.0)
```
The material parameters are read from the tet mesh, and are cached in the constructor. Throws an exception if the provided tet mesh does not consist of materials specifying E, ν.

class HomogeneousNeoHookeanIsotropicMaterial : public IsotropicMaterialWithCompressionResistance   Neo-Hookean material with homogeneous (spatially global) material parameters. Provides functions to compute the energy, gradient, and Hessian of the neo-Hookean material given the three invariants I, II, III.

```
HomogeneousNeoHookeanIsotropicMaterial(double E, double nu,
    int enableCompressionResistance=0, double compressionResistance=0.0)
```
Initializes the Young's modulus `E` and the Poisson ratio `nu` of the material.

class MooneyRivlinIsotropicMaterial : public IsotropicMaterialWithCompressionResistance
Mooney-Rivlin material that supports spatially varying material parameters.

```
MooneyRivlinIsotropicMaterial(TetMesh * tetMesh,
    int enableCompressionResistance=0, double compressionResistance=0.0)
```
The material parameters are read from the tet mesh, and are cached in the constructor. Throws an exception if the provided tet mesh does not consist of Mooney-Rivlin materials. Provides functions to compute the energy, gradient, and Hessian of the Mooney-Rivlin material given the three invariants I, II, III. The implemented Mooney-Rivlin material is described in Section 3.5.5 of [Bow09].

class HomogeneousMooneyRivlinIsotropicMaterial : public IsotropicMaterialWithCompressionResistance Mooney-Rivlin material with homogeneous (spatially global) material parameters. Provides functions to compute the energy, gradient, and Hessian of the Mooney-Rivlin material given the three invariants I, II, III.

```
HomogeneousMooneyRivlinIsotropicMaterial(double mu01, double mu10, double v1,
    int enableCompressionResistance=0, double compressionResistance=0.0)
```
Initializes the Mooney-Rivlin material parameters μ_{01}, μ_{10}, v_1.

2.21   libiglInterface

Wrapper for libigl-related routines. Used by `immersionMesher`.

2.22   lighting

class Lighting   Provides tools to read an OpenGL lighting configuration from a file (using the `configFile` library) and apply it to the current rendering context. Up to eight light sources are supported.

```
Lighting(const char * configFilename)
```
Reads the specified file and extracts lighting information. Throws an exception if `ConfigFile` reports an error (i.e., an unspecified mandatory option) in the file.

```
void LightScene()
```
Enables `GL_LIGHTING` and applies the lighting settings contained in the class.

2.23   ListIO

class ListIO   Named `LoadList` before version 4.0. Utilities for saving and loading lists of integers to and from text files.

```
static int load(const char * filename, int * numListEntries, int ** listEntries,
    int offset=0, bool sortAfterLoad = true)
```
Loads a list of integers from `filename`, writing the size of this list to the `int` pointed to by `numListEntries`, and writing the address of the newly-allocated list of read integers to the `int *` pointed to by `listEntries`. Optionally, each integer has `offset` subtracted from it when read, and the integers can be sorted after reading. Returns 1 if an error occurs, and 0 otherwise.

```
static int save(const char * filename, int numListEntries, int * listEntries, int offset=0)
```

Saves to `filename` a list of `numListEntries` integers pointed to by `listEntries`. Optionally, each integer has `offset` *added* to it before being written. Returns 1 on an error, and 0 otherwise.

2.24   massSpringSystem

Provides a collection of classes to initialize and work with mass-spring systems.

class MassSpringSystem   Implements a mass-spring system, a collection of weighted particles connected by springs.

```
MassSpringSystem(int numParticles, double * masses,
    double * restPositions, int numEdges, int * edges,
    int * edgeGroups, int numMaterialGroups,
    double * groupStiffness, double * groupDamping)
```
Constructs a mass-spring system of `numParticles` particles and `numEdges` springs. `masses` specifies the mass of each particle, and `restPositions` contains three subsequent entries for the three-dimensional coordinates of each particle. `edges` contains $2 \cdot$ `numParticles` particle indices, with each pair defining the endpoints of a spring. `edgeGroups` specifies for each spring the index of which of the `numMaterialGroups` material property groups it belongs to, while `groupStiffness` and `groupDamping` specify the stiffness and damping parameters for each of these groups.

```
MassSpringSystem(int numParticles, double * restPositions,
    int numQuads, int * quads, double surfaceDensity,
    double tensileStiffness, double shearStiffness,
    double bendStiffness, double damping)
```
Constructs a mass-spring system based upon an organization of the `numParticles` particles into `numQuads` quads in a quad mesh. The $4 \cdot$ `numQuads`-entry array `quads` specifies each quad with four subsequent particle indices. Particle rest positions are defined as above, and particle masses are defined based upon `surfaceDensity` and the quad surface areas. The stiffness and damping settings determine the spring properties.

```
MassSpringSystem(int numParticles, double * restPositions,
    MassSpringSystemElementType elementType, int numElements,
    int * elements, double density, double tensileStiffness,
    double damping)
```
Constructs a mass-spring system based upon a tetrahedral mesh or cubic mesh, depending on whether `elementType` is set to `TET` or `CUBE`. The `numParticles` particles are organized into `numElements` tetrahedra or cubes, which are specified in the array `elements` (of size 4 times or 8 times `numElements`) by four or eight subsequent particle indices per element. Particle masses are set according to `density` and the element volumes, while spring properties are set according to the stiffness and damping variables.

```
    void ComputeEdgeInfo(int edgeId, const double * u, double * energy, double f[6], double
K[36])
```
Computes the energy, internal forces and stiffness matrix of a single spring in the system.

class MassSpringSystemFromTetMesh
```
    static int GenerateMassSpringSystem(TetMesh * tetMesh,
    MassSpringSystem ** massSpringSystem, double density,
    double tensileStiffness, double damping)
```
Allocates a new `MassSpringSystem` object based upon the tetrahedral mesh `tetMesh` and the specified material properties, and writes its pointer to `*massSpringSystem`. Returns 1 on error, 0 otherwise.

class MassSpringSystemFromTetMeshConfigFile

```
int GenerateMassSpringSystem(const char * configFilename,
    MassSpringSystem ** massSpringSystem,
    MassSpringSystemTetMeshConfiguration * massSpringSystemTetMeshConfiguration = NULL)
```

Allocates a new `MassSpringSystem` based upon the contents of the file `configFilename`, and writes its pointer to `*massSpringSystem`. If the last argument is not `NULL`, the options read from the config file are written to the pointed-to object. Returns 1 on error, 0 otherwise.

class MassSpringSystemFromObjMesh

```
int GenerateMassSpringSystem(ObjMesh * quadMesh,
    MassSpringSystem ** massSpringSystem, double surfaceDensity,
    double tensileStiffness, double shearStiffness,
    double bendStiffness, double damping)
```

Allocates a new `MassSpringSystem` based upon a `.obj` mesh `quadMesh` and the given material properties, and writes its pointer to `*massSpringSystem`. `quadMesh` must have only quadrilateral faces. Returns 1 on error, 0 otherwise.

class MassSpringSystemObjMeshConfiguration

```
int GenerateMassSpringSystem(const char * configFilename,
    MassSpringSystem ** massSpringSystem,
    MassSpringSystemObjMeshConfiguration * massSpringSystemObjMeshConfiguration = NULL)
```

Allocates a new `MassSpringSystem` based upon the contents of the file `configFilename`, and writes its pointer to `*massSpringSystem`. If the last argument is not `NULL`, the options read from the config file are written to the pointed-to object. Returns 1 on error, 0 otherwise.

class MassSpringSystemFromCubicMesh

```
static int GenerateMassSpringSystem(CubicMesh * CubicMesh,
    MassSpringSystem ** massSpringSystem, double density,
    double tensileStiffness, double damping)
```

Allocates a new `MassSpringSystem` based upon the given `CubicMesh` and material properties, and writes its pointer to `*massSpringSystem`. Returns 1 on error, 0 otherwise.

class MassSpringSystemCubicMeshConfiguration

```
int GenerateMassSpringSystem(const char * configFilename,
    MassSpringSystem ** massSpringSystem,
    MassSpringSystemCubicMeshConfiguration * massSpringSystemCubicMeshConfiguration = NULL)
```

Allocates a new `MassSpringSystem` based upon the contents of the file `configFilename`, and writes its pointer to `*massSpringSystem`. Returns 1 on error, 0 otherwise.

class RenderSprings

```
void Render(MassSpringSystem * massSpringSystem, double * u)
```

Uses `GL_LINES` to draw each of the springs in the mass-spring system pointed to by `massSpringSystem` given the current particle displacements `u`. Edges are color-coded based on which material group they belong to.

2.25   matrixIO

Provides functions to write and read dense matrices to and from a file. The templated functions, where type `real` may be `float` or `double`, operate on `real`-typed matrices stored as column-major arrays. If an error occurs, functions ending in an underscore exit with non-zero status, and functions without an underscore return a non-zero integer.

```
int ReadMatrixFromDisk(const char* filename, int * m, int * n, real ** matrix)
int ReadMatrixFromDisk_(const char* filename, int * m, int * n, real ** matrix)
```
Reads a binary-format matrix from `filename` into a newly-allocated (`malloc`) array, writes the pointer to this array to `*matrix`, and writes the matrix dimensions to `*m` and `*n`. An error occurs if the file reports a different matrix size or if file IO fails.

```
int WriteMatrixToDisk(const char* filename, int m, int n, real * matrix)
int WriteMatrixToDisk_(const char* filename, int m, int n, real * matrix)
```
Writes an `m`-by-`n` matrix stored in the array `matrix` to `filename`. An error occurs if file IO fails.

```
int ReadMatrixFromDiskTextFile(const char* filename, int * m, int * n, real ** matrix)
int WriteMatrixToDiskTextFile(const char* filename, int m, int n, real * matrix)
```
Read and write matrices similarly to `ReadMatrixFromDisk` and `WriteMatrixToDisk`, but encode the `float` or `double` values in text format.

```
int ReadMatrixFromDiskListOfFiles(const char * fileList, int * m, int * n,
    real ** matrix)
```
Concatenates the columns of the matrices from several (binary-format) matrix files into a single matrix. The file of name `fileList` is read, and each line of the file is assumed to be a matrix filename. Each of these matrices must have the same number of rows. The resulting concatenation is written to `*matrix`, and its dimensions are written to `*m` and `*n`.

2.26 matrix

Implements a matrix: a 2D array of real values, together with the commonly defined algebraic operations. The matrix can be rectangular (need not be square). The matrix is stored in the column-major order (same as in LAPACK). Storage is dense (see the `sparseMatrix` library for sparse matrices). The class supports common algebraic operations (addition, multiplication, transposition, etc.), including operator overloading, so you can write expressions such as:

```
A += 0.25 * (B + A) * C;
```

The library can also perform more complex matrix operations such as computing the SVD decomposition, solving linear systems (via LU decomposition), finding eigenvalues and eigenvectors, solving least square systems, and computing matrix inverses, pseudoinverses and exponentials. The code uses BLAS routines to perform matrix addition and multiplication, and LAPACK for SVD, LU, eigenanalysis, least square systems, inverses and pseudoinverses. Expokit [Sid98] is used for matrix exponentiation (disabled by default; enable it in `matrix.h`, `USE_EXPOKIT` macro). The classes are templated for both float and double datatypes. Also included are routines to perform Principal Component Analysis (PCA) on the columns of the matrix.

2.27 mesh

Provides several low-level geometric queries (computing triangle angles and normals, etc.), implements bounding boxes, simple triangle and tet meshes storage classes for geometric purposes (as opposed to `objMesh` and `volumetricMesh` that are used for rendering and simulating purposes), exact predicates based on code from Jonathan Shewchuk [She97] (https://www.cs.cmu.edu/~quake/robust.html), and Tunicate library (http://www.cs.ubc.ca/~rbridson/tunicate/), and an octree for exact queries.

2.28 mesher

Implements Delaunay tetrahedralization of an arbitrary point set in 3D [Bow81] (including efficiently adding new points), quality triangular meshing of distance field level sets [ORY05], and (experimental feature) constrained Delaunay 3D tet meshing with refinement [Si08, SG11]. Mesher examples can be found in `examples/mesher`.

class DelaunayMesher   Implements Delaunay tetrahedralization of an arbitrary point set in 3D, by following the well-known Bowyer's algorithm in 3D [Bow81]. Our implementation makes it possible to efficiently insert new points into the tetrahedralization. We provide an implementation using exact predicates from Shewchuk [She97].

```
bool computeDelaunayTetrahedralization(const std::vector<Vec3d> & points,
    double epsilon = 1e-9)
```
Computes Delaunay tetrahedralization of a set of input points (must have more than 4 points). The parameter `epsilon` is used to check whether input vertices are degenerate.

```
bool addOnePoint(const Vec3d & p)
```
Incrementally adds one point into the tetrahedralization, and updates the Delaunay tetrahedralization.

```
TetMesh * getMesh()
```
Returns the Delaunay tet mesh corresponding to the current set of points.

class IsosurfaceMesher   This class can generate a quality triangle mesh of a user-specified isosurface of a distance field, as described by [ORY05]. The user can control the distance field level set value, and the desired output size of the triangles. The distance field can be signed or unsigned.

```
IsosurfaceMesher(DistanceFieldBase * distanceField)
```
Initializes the mesher to the given distance field.

```
bool compute(double isovalue, double angularBound, double radialBound,
    int numInitialSamples = 20, double epsilon = 1e-6, int maxNumberOfIterations = -1,
    double maxTimeSeconds = -1.0)
```
Computes an isosurface triangular mesh of the distance field level set corresponding to the given `isovalue`. Parameter `angularBound` specifies that all angles in the output triangle mesh must be larger than `angularBound` degrees. Parameter `radialBound` imposes a maximum value for the largest radius of the "surface Delaunay balls" calculated in the internal process of generating the isosurface mesh (see [ORY05]). Parameter `epsilon` are used for the internal DelaunayMesher, as described with the `DelaunayMesher` class. The method will terminate if the number of iterations in the algorithm exceeds `maxNumberOfIterations`. A value of `maxNumberOfIterations` < 0 means no limitation. The method will also terminate if the computation time (in seconds) exceeds `maxTimeSeconds`. A value of `maxTimeSeconds` < 0 means no limitation. Returns true if a timeout occurred, and false otherwise.

```
ObjMesh * getMesh(bool enforceManifold = true)
```
Returns the output isosurface mesh. If the `enforceManifold` is set to true, mesh will be post-processed to be manifold, by closing holes and performing edge collapses. This process is not guaranteed to always produce a manifold mesh.

class TetMesher   Implements a constrained 3D Delaunay tet mesher with refinement, using the methods from [Si08, SG11]. The input is a manifold and self-intersection-free triangle mesh, and the output is a quality tet mesh of the volume enclosed by the input mesh that conforms to the input mesh. This mesh is generated using constrained Delaunay tetrahedralization, and by inserting new vertices ("Steiner vertices") into the mesh as needed to satisfy the quality conditions.
  Note: The tet mesher is experimental. It may generate sliver tets with some inputs.

```
TetMesh * compute(ObjMesh * surfaceMesh, double refinementQuality = 1.1,
    double alpha = 2.0, double minDihedralAngle = 0.0,
    int maxSteinerVertices = -1, double maxTimeSeconds = -1.0)
```
Computes the tet mesh, using constrained 3D Delaunay tetrahedralization, with refinement. Parameter `refinementQuality` >= 1.0 is a scalar; the smaller the value, the more the mesh will be refined. Parameter

`alpha` controls the closeness of the added Steiner vertices. During refinement, it restricts the smallest distance between a new Steiner vertex and the existing vertices to be larger than alpha times the average edge length of the input mesh. Parameter `minDihedralAngle` specifies a minimal tet dihedral angle in the output mesh. Parameter `maxSteinerVertices` gives an upper bound on the number of Steiner vertices that can be added during the refinement process. Parameter `maxSteinerVertices` < 0 means no limitation. Parameter `maxTimeSeconds` terminates the method if the computation time (in seconds) is exceeded. Parameter `maxTimeSeconds` < 0 means no limitation. Returns the computed tet mesh.

2.29   minivector

class Vec2d   Provides basic functionality for two-dimensional vectors with `double`-valued coordinates. Overloads addition and subtraction of vectors, scalar (pre- and post-)multiplication and division, equality testing, `[]` member access, and `<<` output with `ostream`.

```
Vec2d()
Vec2d(double x, double y)
Vec2d(double entry)
```
Does no initialization, sets the two coordinates to `x` and `y`, or sets both to `entry`.

```
friend double dot(const Vec2d & vec1, const Vec2d & vec2)
```
Returns the dot product of `vec1` and `vec2`.

```
friend Vec2d norm(const Vec2d & vec1)
```
Returns a normalized copy of `vec1`.

class Vec3d   Provides three-dimensional vector functionality with `double`-valued coordinates. Overloads addition and subtraction of vectors, scalar (pre- and post-)multiplication and division, equality testing, member access, and `<<` output with `ostream`.

```
Vec3d()
Vec3d(double x, double y, double z)
Vec3d(double entry)
Vec3d(const double * vec)
Vec3d(const Vec3d & vec)
```
Does no initialization, sets the three coordinates, sets all three coordinates to `entry`, sets coordinates from an array, or copies from another `Vec3d`.

```
friend double dot(const Vec3d & vec1, const Vec3d & vec2)
```
Returns the dot product of `vec1` and `vec2`.

```
friend Vec3d cross(const Vec3d & vec1, const Vec3d & vec2)
```
Returns the cross product of `vec1` and `vec2`.

```
friend Vec3d norm(const Vec3d & vec1)
```
Returns a normalized copy of `vec1`.

class Mat3d   Implements a three-dimensional matrix. Overloads operators for addition and subtraction of matrices, scalar (pre- and post-) multiplication, and `<<` output for `ofstream`. The `[]` operator can be used to access the row vectors, whose elements can be accessed with a second `[]`.

```
Mat3d()
Mat3d(double x0, double x1, double x2, ..., double x8)
Mat3d(const double mat[9])
Mat3d(const Vec3d rows[3])
```

```
Mat3d(const Vec3d & row0, const Vec3d & row1, const Vec3d & row2)
Mat3d(double diag)
```
Does no initialization, initializes all 9 entries (in row-major order) either individually or with an array, sets the three rows either from an array or individually, or initializes a diagonal matrix.

```
void set(double x0, double x1, double x2, ..., double x8)
void set(double value)
```
Set the values of all 9 matrix entries in row-major order, or set all entries to `value`.

```
friend Mat3d tensorProduct(const Vec3d & vec1, const Vec3d & vec2)
```
Returns the tensor product of the two vectors. If `vec1` and `vec2` are interpreted as column vectors, column `i` of the output is `vec1` scalar-multiplied by `vec2[i]`.

```
friend Mat3d inv(const Mat3d &)
```
Returns the inverse of the input matrix. Performs no checking to ensure the inverse exists.

```
friend double det(const Mat3d & mat)
```
Returns the determinant of `mat`.

```
friend Mat3d trans(const Mat3d & mat)
```
Returns the transpose of `mat`.

```
void convertToArray(double * array)
```
Writes the matrix entries to the array in row-major order.

```
friend void eigen_sym(Mat3d & M, Vec3d & eig_val, Vec3d eig_vec[3])
```
Calculates the eigenvalues and eigenvectors of matrix M and writes the eigenvalues to `eig_val` and the eigenvectors to `eig_vec`. This routine calls a public domain routine `eigen_decomposition`, available in `eig3.h` and `eig3.cpp`. These two files were downloaded from: `http://barnesc.blogspot.com/2007/02/eigenvectors-of-3x3-symmetric-matrix.html`. The website states that the code is public domain, and that it was obtained from the public-domain code "JAMA : A Java Matrix Package": `http://math.nist.gov/javanumerics/jama/`.

```
friend void svd(Mat3d & M, Mat3d & U, Vec3d & Sigma, Mat3d & V,
    double singularValue_eps=1e-8, int modifiedSVD=0)
```
Given a 3x3 matrix M, decomposes it using the singular value decomposition so that $M = U\Sigma V^T$, where U is a 3x3 rotation, V is 3x3 orthonormal ($V^T V = V V^T = I$), and Σ is a diagonal matrix with non-negative entries, in descending order, $\Sigma[0] \geq \Sigma[1] \geq \Sigma[2] \geq 0$. Note that $det(V)$ may be 1 (rotation) or -1 (reflection). Parameter `singularValue_eps` is a threshold to determine when a singular value is deemed zero, and special handling is then invoked to improve numerical robustness. Parameter `modifiedSVD` controls whether SVD is standard or modified. When standard (`modifiedSVD=0`), SVD is performed as described above. When modified (`modifiedSVD=1`), SVD is modified so that it has the following properties (useful in solid mechanics applications): (1) not just the determinant of U, but also the determinant of V is 1 (i.e., both U and V are rotations, not reflections), and (2) the smallest diagonal value $\Sigma[2]$ may be negative, and we have $\Sigma[0] \geq \Sigma[1] \geq abs(\Sigma[2]) \geq 0$.

2.30   modalMatrix

Provides storage for a modal matrix for use in model reduction. Computes world-coordinate vertex positions from the reduced coordinates, using the equation $u = Uq$. The computation is performed on the CPU (using BLAS).

class ModalMatrix   Performs the functionality described above.

```
void AssembleVector(double * q, double * u)
```
Computes world-coordinate vertex positions u using the equation $u = Uq$. The multiplication is performed using the BLAS/LAPACK library, on the CPU.

```
void AssembleMatrix(int numColumns, double * qMatrix, double * uMatrix)
```
Computes several deformations vectors (columns of a matrix) at once.

```
inline void AssembleSingleVertex(int vertex, double * q,
    double * ux, double * uy, double * uz)
```
Computes the displacement of a single vertex.

2.31   objMesh

Provides utilities to load/save `.obj` meshes (including mtl files), access/modify the geometry of a mesh, and render meshes using OpenGL. The mesh storage structure is as follows: there are three global arrays, storing positions of mesh vertices, texture coordinates and normals. Nested classes `ObjMesh::Group`, `ObjMesh::Face` and `ObjMesh::Vertex` then represent the mesh geometry by giving integer indices into these three global arrays, for each vertex on every face. This means that a vertex that is shared by two (or more) faces may have different normals or texture coordinates in the two faces.

class ObjMesh::Vertex   Constructed by the `ObjMesh` constructor. Stores the information about a particular vertex in the mesh.

```
unsigned int getPositionIndex()
```
Returns the index of this vertex's `Vec3d` position vector in the global array of position vectors in the `ObjMesh`. The position vector may be accessed with `ObjMesh::getPosition`.

```
bool hasNormal()
bool hasTextureCoordinate()
```
Returns whether the vertex has a normal vector or a texture coordinate vector. Must be called before requesting the normal vector or texture coordinate indices with the functions below.

```
unsigned int getNormalIndex()
unsigned int getTextureCoordinateIndex()
```
Returns the index of this vertex's `Vec3d` normal vector or texture coordinate vector in the global array of normal vectors or texture coordinate vectors in the `ObjMesh`. These vectors may be accessed with the `ObjMesh::getNormal` and `ObjMesh::getTextureCoordinate` functions, respectively. Aborts if the vertex has no normal or texture coordinate.

```
void setPositionIndex(unsigned int positionIndex)
void setNormalIndex(unsigned int normalIndex)
void setTextureCoordinateIndex(unsigned int textureCoordinateIndex)
```
Sets the position index, normal index, or texture coordinate index for this vertex.

class ObjMesh::Material   Constructed by the `ObjMesh` constructor. Stores material properties settings from an `.obj` file.

```
Vec3d getKa()
Vec3d getKd()
Vec3d getKs()
```
Returns a vector containing the RGB components of the ambient, diffuse, or specular lighting coefficients, respectively.

```
double getShininess()
double getAlpha()
```
Returns the shininess value for specular reflections or the alpha (transparency) value.

```
void setAlpha(double alpha)
```
Sets the alpha (transparency) value of the material.

class ObjMesh::Face Constructed by the `ObjMesh` constructor. Stores a `vector` of `ObjMesh::Vertex` objects that constitute a face of the model.

```
void addVertex(const Vertex& v)
```
Adds vertex `v` to the end of the list of vertices constituting the face.

```
size_t getNumVertices()
```
Returns the number of vertices in the face.

```
const Vertex & getVertex(unsigned int vertex)
Vertex * getVertexHandle(unsigned int vertex)
```
Returns a reference of or a pointer to the face vertex at index `vertex`.

```
bool hasFaceNormal()
```
Returns whether the face has a precomputed face normal.

```
void setFaceNormal(Vec3d & normal)
```
Sets the face normal of the face.

```
Vec3d getFaceNormal()
```
Returns the precomputed face normal, asserting that it has previously been set.

class ObjMesh::Group Constructed by the `ObjMesh` constructor. Stores a `vector` of faces constituting a group in the `.obj` file.

```
void addFace(const Face& face)
```
Adds a face to the `vector` of faces.

```
size_t getNumFaces()
```
Returns the number of faces in the group.

```
const Face & getFace(unsigned int face)
Face getFaceHandle(unsigned int face)
```
Returns a reference of or a pointer to the face at index `face`.

```
std::string getName()
```
Returns the name of the group.

```
unsigned int getMaterialIndex()
void setMaterialIndex(unsigned int index)
```
Gets or sets the material index for this group. The index refers to the global array of materials in the `ObjMesh`.

```
void removeFace(unsigned int i)
```
Removes the face at index `i`.

class ObjMesh Reads and stores the model information from an `.obj` file.

    ObjMesh(const std::string& filename, int verbose=1)
Initializes the object based upon the contents of the `.obj` file `filename`. Throws an `ObjMeshException` if any failure occurs during the process.

    ObjMesh()
Initializes an empty model.

    std::string filename()
Returns the name of the `.obj` file used to initialize the object. Returns an empty string if the object was not constructed from a file.

    void printInfo()
Prints detailed information about the model.

    bool isTriangularMesh()
    bool isQuadrilateralMesh()
Returns whether the mesh is triangular or quadrilateral.

    void triangulate()
Subdivides faces so that the mesh consists only of triangular faces.

    void getBoundingBox(double expansionRatio, Vec3d * bmin, Vec3d * bmax)
Return a bounding box of the mesh. The tightest fitting box is scaled by `expansionRatio`. For `expansionRatio`=1, one obtains the tight-fitting bounding box.

    Vec3d getPosition(int vertexIndex)
    Vec3d getTextureCoordinate(int textureCoordinateIndex)
    Vec3d getNormal(int normalIndex)
Returns the vertex position, texture coordinate, or normal vector held at index `vertexIndex`, `textureCoordinateIndex`, `normalIndex`, in the global array of vertex position, texture coordinate, or normal vectors.

    Vec3d getPosition(Vertex & vertex)
    Vec3d getTextureCoordinate(Vertex & vertex)
    Vec3d getNormal(Vertex & vertex)
Returns the vertex position, texture coordinate, or normal vector for the given vertex.

    void setPosition(int vertexIndex, Vec3d & position)
    void setTextureCoordinate(int textureCoordinateIndex, Vec3d & textureCoordinate)
    void setNormal(int normalIndex, Vec3d & normal)
    void setPosition(Vertex & vertex, Vec3d & position)
    void setTextureCoordinate(Vertex & vertex, Vec3d & textureCoordinate)
    void setNormal(Vertex & vertex, Vec3d & normal)
Sets the vertex position, texture coordinate or normal in the array of global vertex positions, texture coordinates or normals.

    const Material & getMaterial(unsigned int index)
    Material * getMaterialHandle(unsigned int index)
Returns a reference of or a pointer to the material at index `index` in the global array of materials.

    size_t getNum...()
    void add...(...)

For each of `Vertices`, `Faces`, `Normals`, `TextureCoordinates`, `Groups`, or `Materials`, returns the number of them in the object, or adds a new one of them to the object.

```
void buildFaceNormals()
```
Calculates the geometric normals of each face (assuming counter-clockwise winding) and caches them in the `ObjMesh::Face` objects. This must be done before using any of the vertex-normal generation functions below.

```
void buildVertexFaceNeighbors()
void clearVertexFaceNeighbors()
```
(Re-)Builds or clears an internally-stored structure for easy lookup of faces neighboring a given vertex. This must be done before using any of the vertex-normal generation functions below.

```
void buildVertexNormals(double angle)
```
The functions `buildFaceNormals` and `buildVertexFaceNeighbors` will be called by this function if not yet called. Sets the vertex normals of the mesh to create a smooth-looking surface except for at angles exceeding `angle`.

```
void buildVertexNormalsFancy(double angle)
```
The functions `buildFaceNormals` and `buildVertexFaceNeighbors` must be called by this function if not yet called. Uses a slower but more precise method to set the vertex normals of the mesh to create a smooth-looking surface except for at angles exceeding `angle`.

```
void save(const std::string & filename, int outputMaterials = 0,
    fileFormatType fileFormat = BY_EXT, int verbose=1)
```
Saves the mesh to file `filename`, in `obj` format or a binary `objb` format for smaller files and faster loading times. Optionally it writes the mesh materials to `filename`.mtl. `fileFormat` can be ASCII (.obj format), BINARY (.objb format) or BY_EXT. If it is BY_EXT, the function chooses the file format based on the suffix in `filename`.

2.32 objMeshGPUDeformer

Provides routine to deform triangle mesh vertices using a GPU. Two deformer types are supported: (1) interpolating from a volumetric mesh deformation to an embedded triangle mesh, and (2) computing the triangle mesh deformations in model reduction as $u = Uq$, where U is the modal matrix and q are the reduced coordinates.

class ObjMeshRender   Renders an `ObjMesh` object via OpenGL.

```
ObjMeshRender(ObjMesh * mesh)
```
Initializes the object based upon the mesh `mesh`.

```
void render(int geometryMode, int renderMode)
unsigned int createDisplayList(int geometryMode, int renderMode)
```
Renders the object, or returns the integer identifier of an OpenGL display list to render the object. `geometryMode` is a bit field of options, controlling whether to render faces, edges, vertices, and/or normals. `renderMode` is a bit field of options to control rendering settings such as whether to use flat or smooth shading, whether to render textures, and/or whether to use colors or to set `glMaterial` material properties. See `objMeshRender.h` for more detailed info.

```
void renderSpecifiedVertices(int * specifiedVertices, int numSpecifiedVertices)
```
Renders the `numSpecifiedVertices` vertices whose indices are specified in the array `specifiedVertices`.

```
void renderGroupEdges(const char * groupName)
```

Renders all the edges in the group of name `groupName` in the mesh.

`void loadTextures(int textureMode)`
Loads any textures specified in the material file of the mesh used to construct the `ObjMesh` mesh. `textureMode` specifies the color blend mode and mip-mapping settings for the textures; see `ObjMeshRender::Texture::loadTexture`.

`static unsigned char * loadPPM(std::string filename, int * width, int * height)`
Loads a binary-format RGB .ppm file (type P6), writes the dimensions to `*width` and `*height`, and returns a pointer to a newly-allocated (`new`) array of size $3 \cdot \text{width} \cdot \text{height}$ of the components of the pixels of the image. The order of pixels in the output is bottom row to top row. Returns `NULL` upon failure.

class ObjMeshRender::Texture   Stores the information for an OpenGL texture.

`Texture()`
Initializes the object to store no texture.

`~Texture()`
If the object stores a texture, calls `glDeleteTextures` on it.

`void loadTexture(std::string fullPath, int textureMode)`
Loads the texture stored in the binary-format `.ppm` file at `fullPath` (path from location of executable to texture file) and writes it to an OpenGL texture. `textureMode`, a bit field of options OBJMESHRENDER_GL_REPLACE or OBJMESHRENDER_GL_MODULATE for blend mode and OBJMESHRENDER_GL_NOMIPMAP or OBJMESHRENDER_GL_USEMIPMAP for whether to use mip-mapping, indicates the settings for this texture.

`bool hasTexture()`
Returns whether this object has loaded a texture.

`unsigned int texture()`
Returns OpenGL's identifier for the texture this object holds. Asserts that a texture has been loaded.

`int textureMode()`
Returns the texture mode as passed to this object in `loadTexture`.

2.33   openGLHelper

Provides a handful of functions for assorted tasks such as saving a screenshot of the OpenGL window, rendering a unit cube, or arrows.

2.34   performanceCounter

class PerformanceCounter   Used to measure the execution time of a block of code. Construct the timer before a block of code to begin timing, stop it explicitly after the block executes, and then read the elapsed time at roughly microsecond accuracy.

`PerformanceCounter()`
Starts the counter at the current system time. If needed, the start time can be explicitly reset later (see below).

`void StartCounter()`
Restarts the counter at the current system time.

`void StopCounter()`

Stops the counter at the current system time.

```
double GetElapsedTime()
```
Returns the elapsed time between when the timer was last started and when it was last stopped. Note that one cannot use a single timer to measure the total running time of two blocks of code by starting and stopping it twice.

2.35   polarDecomposition

Provides tools to calculate the polar decomposition and the gradient of the polar decomposition of a 3×3 matrix. The code in the `PolarDecomposition` class is an adapted version of the code published by Graphics Gems IV, Ken Shoemake, 1993, "Polar Decomposition of 3x3 matrix in 4x4, M = QS.", available at: `http://tog.acm.org/GraphicsGems/`. The website states that "Using the code is permitted in any program, product, or library, non-commercial or commercial." For details, see the header to `polarDecomposition.h`.

class PolarDecomposition
```
static double Compute(const double * M,
    double * Q, double * S, double tol = 1.0e-6)
```
Calculates the polar decomposition of input matrix `M` into matrices `Q` and `S` (within an accuracy threshold `tol`) such that `M = QS`, `Q` is a 3×3 orthogonal matrix, and `S` is a 3×3 symmetric matrix. All matrices are encoded as arrays in row-major format, and `Q` and `S` are assumed to be already allocated.

class PolarDecompositionGradient
```
static void Compute(const double * M, const double * Q,
    const double * S, const double * MDot, double * omega,
    double * QDot, double * SDot, const double * MDotDot=NULL,
    double * omegaDot=NULL, double * QDotDot = NULL)
```
Takes as input matrices `M`, `Q`, `S` (such that `M = QS` is the polar decomposition), 3×3 derivative of `M` `MDot`, and optionally the 3×3 second derivative `MDotDot`. Calculates 3×3 derivatives `QDot` and `SDot`, and 3-vector of rotational velocity `omega`, along with second derivatives `QDotDot` and `omegaDot` if `MDotDot` is provided and the output pointers are not `NULL`. Again, all matrices are stored as row-major arrays, and all non-optional output matrix arrays must be pre-allocated.

2.36   quaternion

Implements quaternions and the common algebraic operations on quaternions (including operator overloading). The class is templated: you can use either float or double precision. Supports using quaternions to represent/manipulate rotations.

2.37   reducedElasticForceModel

This library provides several concrete (non-abstract) classes to be used by the `integratorDense` library, such as a reduced StVK wrapper (via the `reducedStvk` library), reduced mass spring system and linearized versions of such models. This library parallels `elasticForceModel` in spirit; see documentation for `elasticForceModel`. This way, for example, the `reducedStvk` library can be used together with `integratorDense`.

2.38   reducedForceModel

This library provides the abstract base class for a force model used in the `integratorDense` library, i.e., a "black-box" function $q \mapsto U^T f_{\text{int}}(Uq)$ and its gradient in

$$\tilde{M}\ddot{q} + \left(\alpha\tilde{M} + \beta U^T K(Uq)U + \tilde{D}\right)\dot{q} + U^T f_{\text{int}}(Uq) = U^T f_{\text{ext}}.$$

Any dynamical system described by such a differential equation can then be timestepped by the `integratorDense` library, by providing an implementation of $U^T f_{\text{int}}(Uq)$ and its gradient, in a class derived from `reducedForceModel`. This library parallels `forceModel` in spirit; see documentation for `forceModel`.

class ReducedForceModel   Abstract base class implementing the functionality described above.

2.39   rigidBodyDynamics

Implement 6-DOF rigid dynamics of a single rigid body, as explained in [BW01], under any specified (time-varying) external forces and torques. Arbitrary tensors of inertia are supported. For rigid objects where the inertia tensor in the world coordinate system is diagonal, use `RigidBody`. For the general case (non-diagonal inertia tensor in the world coordinate system), use `RigidBody_GeneralTensor`. The solution is computed by numerically timestepping the ordinary differential equations of rigid body motion, derived from the Newton's 2nd law, and conservation of linear momentum and angular momentum. For example, ballistic motion can be simulated if gravity is used as the external force. Objects bouncing off the ground/impacting other objects can be simulated if you combine this library with a collision detection algorithm that provides the contact external forces. The code supports explicit Euler integration and symplectic Euler integration.

class RigidBody   A rigid body with a diagonal inertia tensor in the rest configuration.

class RigidBody_GeneralTensor   Rigid body with a general (potentially non-diagonal) tensor in the rest configuration.

2.40   reducedStvk

Implements the model-reduced St. Venant Kirchhoff deformable model. It provides classes to evaluate reduced internal forces, stiffness matrix, and Hessian tensor (derivative of stiffness matrix) for a given volumetric mesh, its deformation basis matrix, and the reduced coordinates. Each element of the force vector is a cubic polynomial in the reduced coordinates `q`. The force vector and reduced coordinates are vectors of length `r`, where `r` is the dimension of the reduced basis(see [Bar07]). Optionally, the computation of reduced internal forces can be multi-threaded (using the Intel TBB library).

class StVKReducedInternalForces   Precomputes the coefficients of the cubic polynomials, and efficiently evaluates the reduced internal forces for the StVK material, for any reduced configuration q.

```
StVKReducedInternalForces(int r, double * U, VolumetricMesh * volumetricMesh,
    StVKElementABCD * precomputedABCDIntegrals, int initOnly=0,
    bool addGravity=false, double g=9.81, int verbose=1)
```
Precomputes cubic reduced internal force polynomials, for the given mesh `volumetricMesh`, deformation basis `U`, and precalculated mesh information `precomputedABCDIntegrals` (usually generated by `StVKElementABCDLoader` given the mesh).

```
void Evaluate(double * q, double * fq)
```
Given the reduced coordinates `q`, computes reduced internal forces and writes them to the pre-allocated array `fq`.

```
void SetGravity(bool addGravity, double g,
    VolumetricMesh * volumetricMesh_=NULL, double * U_=NULL)
```
Sets whether to apply the gravitational force `g`.

```
void Scale(double scalingFactor)
```
Scales the stiffness of the model. Consequently, the frequency spectrum will be scaled linearly by the square root of `scalingFactor`.

```
StVKReducedInternalForces * ShallowClone()
```
Makes shallow copies of all pointers, except those initialized by the function `InitBuffers`. This ensures thread safety if the user wants to evaluate two or more identical models (i.e., two copies of an object) in parallel. Note that this routine is not needed if evaluating the reduced internal forces for a single model.

2.41   sceneObject

Allows for easy rendering of `.obj` meshes, either in their original state or under user-specified deformations. File-loading and mesh-rendering functionality is provided by the `obj` library. All indices taken or returned by the classes below are zero-indexed.

class SceneObject   Stores the model from an `.obj` file and provides tools to render the entire model or its constituent features.

```
SceneObject(const char * filename)
```
Constructs the object from the .obj file at `filename`. Calls `exit` upon failure.

```
virtual void Render()
```
Renders the entire model with an `ObjMeshRender` object.

```
virtual void RenderVertices()
virtual void RenderEdges()
virtual void RenderNormals()
```
Render the specified components of the model using `ObjMeshRender`.

```
virtual void RenderVertex(int vertex)
```
Renders the vertex at index `vertex` in the mesh.

```
void HighlightVertex(int i)
```
Renders vertex `i` with a large green dot.

```
void BuildDisplayList()
```
Builds a display list for the model to speed up rendering. If a list has been built, it is automatically used by the `Render` function above.

```
void PurgeDisplayList()
```
Deletes the display list, if there currently is one.

```
ObjMesh * GetMesh()
```
Returns a pointer to the underlying `ObjMesh` object from the `obj` library, for more direct access to the mesh.

```
void ComputeMeshRadius(Vec3d & centroid, double * radius)
```
Writes to `*radius` the smallest radius of a sphere centered at `centroid` that contains the entire mesh.

```
void ComputeMeshGeometricParameters(Vec3d * centroid, double * radius)
```
Writes the coordinates of the center of the mesh (the average of its vertices) to `*centroid`, and the smallest radius of a sphere around this center which contains the entire mesh to `*radius`.

```
void ExportMeshGeometry(int * numVertices, double ** vertices,
    int * numTriangles, int ** triangles)
```

Writes to *numVertices and *numTriangles the number of vertices and triangles in the mesh, and writes to *vertices and *triangles pointers to newly-allocated (malloc) arrays of the three coordinates of each vertex and the three vertex indices of each triangle, respectively.

```
virtual int GetClosestVertex(Vec3d & queryPos, double * distance=NULL,
    double * auxVertexBuffer = NULL)
```

Returns the index of the mesh vertex closest to the specified query location. The auxVertexBuffer option is only need for derived classes; you may pass NULL when using GetClosestVertex with SceneObject. To find the nearest vertex in a mesh that has been deformed, you need to use the SceneObjectDeformable subclass.

```
int SetUpTextures(LightingModulationType lightingModulation, MipmapType mipmap)
```

Initializes any textures specified by the .obj file, and enables texture rendering. Calling this function is mandatory for texture mapping to work. Must be called after OpenGL has been initialized. The lightingModulation parameter selects either MODULATE or REPLACE for the texture blend mode, and the mipmap parameter selects either USEMIPMAP or NOMIPMAP for whether to use mip-mapping. Returns 0 iff successful.

```
bool hasTextures()
```

After SetUpTextures has been called, returns whether the model has textures.

```
void EnableTextures()
void DisableTextures()
```

Enables or disables textures for rendering. Does not check whether textures have been set up.

```
void BuildFaceNormals()
```

Calculates the normals of the mesh faces.

```
void BuildNeighboringStructure()
```

Must be called before any of the three vertex normal methods below. Creates a data structure to optimize finding neighboring faces of mesh vertices.

```
void BuildVertexNormals(double thresholdAngle=85.0)
```

Must be called after BuildNeighboringStructure. Calculates smooth normals for each vertex by averaging the face normals of each incident triangle. The thresholdAngle parameter preserves sharp angles in the mesh; for a given vertex, if a triangle incident to it has angle greater than thresholdAngle between its normal and the normal of the first triangle incident to the vertex, then this triangle does not contribute to or use the smooth normal at the vertex.

```
void BuildNormals(double thresholdAngle=85.0)
```

Must be called after BuildNeighboringStructure. Builds face and vertex normals for the mesh.

```
void BuildNormalsFancy(double thresholdAngle=85.0)
```

Must be called after BuildNeighboringStructure. Sets mesh normals, handling sharp edges in a more precise way.

```
void ShowPointLabels(int k, int l)
```

Renders vertex indices for vertices with indices k through l. The values are rendered as one-indexed.

```
void ShowPointLabels()
```

Renders all vertex indices.

class SceneObjectWithRestPosition : public SceneObject   Adds to SceneObject a record of the rest position of the model, i.e. the initial position of each vertex.

```
SceneObjectWithRestPosition(const char * filename)
```
Loads the `.obj` mesh at `filename`, and saves a record of the rest position of the mesh.

```
void GetVertexRestPositions(float * buffer)
void GetVertexRestPositions(double * buffer)
```
Writes the rest position coordinates of each vertex in the mesh to the pre-allocated array `buffer`, in either `float` or `double` format.

```
double * GetVertexRestPositions()
```
Provides direct access to the internal array of rest positions.

class SceneObjectDeformable : public SceneObjectWithRestPosition Adds the ability to deform the model to a given displacement from the rest position, so it may be rendered in the deformed position, and to reset the model to an undeformed state.

```
SceneObjectDeformable(const char * filenameOBJ)
```
Initializes the model from the `.obj` file at `filenameOBJ`.

```
void GetSingleVertexRestPosition(int vertex, double * x, double * y, double * z)
void SetSingleVertexRestPosition(int vertex, double x, double y, double z)
```
Gets or sets the rest position of the vertex at index `vertex` in the mesh.

```
void GetSingleVertexPositionFromBuffer(int vertex,
    double * x, double * y, double * z)
```
Returns the current model position of the vertex at index `vertex`.

```
void ResetDeformationToRest()
```
Resets all vertices to their rest positions.

```
void SetVertexDeformations(double * u)
void SetVertexDeformations(float * u)
```
For array `u` which specifies displacements for the coordinates of each vertex in the model, sets the model to be displaced by `u` from its rest position.

```
void AddVertexDeformations(double * u)
```
Adds displacements `u` of the coordinates of each model vertex to the current positions of the vertices.

```
virtual void SetLighting(Lighting * lighting)
```
Sets the scene lighting using the specified `Lighting` object.

2.42   sceneObjectReduced

Allows for easy rendering of meshes simulated using model reduction. Library supports computing world-coordinate vertex displacements from reduced coordinates. This computation can be performed either on the CPU or the GPU (fragment shader). Note that all classes have a "6DOF" version which permits setting a rigid transformation to be applied to the mesh, useful for simulating deformable objects undergoing rigid body motion.

class SceneObjectReduced Abstract base class for rendering meshes that are simulated using model reduction. The derived classes compute the world-coordinate vertex displacements u via the formula $u = Uq$, where U is the modal matrix, and q are the reduced coordinates.

class SceneObjectReducedCPU   Computes the world-coordinate vertex displacements by multiplying $u = Uq$ using the CPU.

class SceneObjectReducedGPU   Computes the world-coordinate vertex displacements using the GPU, using the class `objMeshGPUDeformer`. The multiplication $u = Uq$ is performed in a fragment shader. If needed, user can read-back the result u to the CPU.

2.43   shapeEdit

Implements the As-Rigid-As-Possible (ARAP) method for fast geometric shape deformation, as described in the following reference:

```
Olga Sorkine and Marc Alexa:
As-Rigid-As-Possible Surface Modeling,
Eurographics Symposium on Geometry Processing, 2007.
```

2.44   sparseMatrix

Provides tools for efficiently storing and running calculations with matrices that have few non-zero elements. These matrices do not need to be square. Rows and columns are zero-indexed.

class SparseMatrixOutline   Used to declare which positions will be non-zero in a `SparseMatrix`. This class cannot perform matrix arithmetic, but non-zero entries can be freely added. Then, the class is used to construct a `SparseMatrix` object, which performs arithmetic but whose zero positions cannot be changed. Contained in the same file as `SparseMatrix`.

```
    SparseMatrixOutline(int n)
    SparseMatrixOutline(int n, double diagonal)
    SparseMatrixOutline(int n, double * diagonal)
```
Sets the number of rows in the matrix, and optionally a starting value for all the diagonal entries or an array of starting values for the diagonals.

```
    SparseMatrixOutline(const char * filename, int expand=1)
```
Constructs the outline from a file. For `expand` greater than 1, expands each element to a diagonal block of size `expand` by `expand`. This is useful, for example, with mass matrices.

```
    void AddEntry(int i, int j, double value=0.0)
    void AddBlock3x3Entry(int i, int j, double * matrix3x3)
    void AddBlockMatrix(int i, int j, const SparseMatrix * block, double scalarFactor=1.0)
```
Adds a value (or a matrix of values) starting at row `i`, column `j`. "Add" in this context means addition, not insertion; if the position in the matrix is already non-zero, this new value is added to the existing value. `scalarFactor` is multiplied with the block matrix values before they are added.

```
    double GetEntry(int i, int j)
```
Gets the matrix value at row `i`, column `j`, or zero if the position has not been set.

class SparseMatrix   A matrix storage class optimized for the case of the matrix having few non-zero elements, i.e., a sparse matrix. The class uses row-based sparse storage. Each matrix row is stored as a list of non-zero column positions and the corresponding entries. The class can perform matrix arithmetic, as well as many other operations such as assigning submatrices or removing matrix rows and columns. It is constructed by first specifying the locations of non-zero entries using the `SparseMatrixOutline` class. Once this is finalized, one can initialize `SparseMatrix`, which writes the row elements into a linear array for fast subsequent access. The sparsity pattern cannot be changed by `SparseMatrix`.

```
SparseMatrix(SparseMatrixOutline * sparseMatrixOutline)
SparseMatrix(const SparseMatrix & source)
SparseMatrix(const char * filename)
```
Initializes the matrix values from a `SparseMatrixOutline`, another `SparseMatrix`, or a file.

```
int Save(const char * filename, int oneIndexed=0)
```
Saves the matrix to disk. Returns 1 on error, 0 otherwise.

```
void SetEntry(int row, int j, double value)
void AddEntry(int row, int j, double value)
void GetEntry(int row, int j)
```
Sets, adds to, or gets the value in the jth *non-zero* element of row `row` of the matrix.

```
int GetRowLength(int row)
```
Returns the number of non-zero elements in the row.

```
int GetColumnIndex(int row, int j)
```
Returns the column index of the jth *non-zero* element in row `row`.

```
SparseMatrix operator+(const SparseMatrix & mat2)
SparseMatrix operator-(const SparseMatrix & mat2)
```
Returns a `SparseMatrix` equal to the sum of this matrix and `SparseMatrix mat2`. These two matrices must have the same pattern of non-zero entries. Same for subtraction.

```
void MultiplyMatrix(int numDenseRows, int numDenseColumns, const double * denseMatrix,
    double * result)
```
Multiplies this `SparseMatrix` matrix with a dense matrix `denseMatrix` of dimensions `numDenseRows` $\times$ `numDenseColumns`. The dense matrices are stored as column-major arrays of doubles, and the output matrix `result` is assumed to be previously allocated.

2.45   sparseSolver

Offers a common interface for classes which solve linear systems, and implements it for Conjugate Gradients, PARDISO and SPOOLES.

class LinearSolver   Serves as an abstract base class for solvers of linear systems.

```
virtual int SolveLinearSystem(double * x, const double * rhs) = 0
```
Implementations solve the linear system $Ax = $ `rhs`, where the matrix A has been specified earlier via some implementation-specific method.

class CGSolver : public LinearSolver   Solves the linear system $Ax = b$ using Conjugate Gradients. The implementation closely follows the reference [She94]. The matrix A must be symmetric positive-definite. Here, x is represented by a dense array of doubles, and A is represented either by a `SparseMatrix` or a user-provided function that calculates Ax for an input x.

```
CGSolver(SparseMatrix * A)
CGSolver(int n, blackBoxProductType callBackFunction, void * data)
```
Initializes the solver from a `SparseMatrix` or a user-provided function as described above. `blackBoxProductType` is a function that returns `void` and accepts `const void *` (pointer to matrix A), `const double *` (pointer to vector x), and `double *` (pointer to product Ax).

```
int SolveLinearSystemWithoutPreconditioner(double * x, const double * b,
    double eps=1e-6, int maxIter=1000, int verbose=0)
```
Solves the equation $Ax = b$ for an input b and the current value of the matrix A. Output is given on the pre-allocated vector x, which should be initialized to an initial guess for the solution. The function terminates when the number of iterations exceeds the input `maxIter` or the error falls below the input `eps`; the return value is the number of iterations taken, multiplied by -1 if the system did not converge.

```
int SolveLinearSystemWithJacobiPreconditioner(double * x, const double * b,
    double eps, int maxIter, int verbose=0)
```
Solves the equation $Ax = b$ using Jacobi preconditioning. Note that this will fail if the diagonal of A contains any zero entries.

```
virtual int SolveLinearSystem(double * x, const double * b)
```
Implements the virtual function from `LinearSolver` by calling `SolveLinearSystemWithJacobiPreconditioner` with the default parameters, eps=1e-6, maxIter=1000, verbose=0.

class PardisoSolver : public LinearSolver Solves the linear system $Ax = b$ using the PARDISO solver. Matrix A must be symmetric and invertible.

```
PardisoSolver(const SparseMatrix * A, int numThreads = 1,
    matrixType mtype = REAL_SYM_INDEFINITE, reorderingType rtype = NESTED_DISSECTION,
    int directIterative=0, int verbose=0)
```
Initializes the solver based upon symmetric matrix `A`. The matrix used for solving the linear system may be changed later, but it must have the same topology as `A`. `numThreads` threads are used for the calculation, `mtype`, `rtype` and `directIterative` declare the type of the matrix (positive definite, etc.), the reordering method of the solver, and whether to use the direct iterative solving method.

```
int ComputeCholeskyDecomposition(const SparseMatrix * A)
```
Performs complete Cholesky factorization on `A`. The factorization will be used in subsequent solves using `SolveLinearSystem`. Returns the error status of the `pardiso` function.

```
virtual int SolveLinearSystem(double * x, const double * rhs)
```
Solves the linear system `Ax = rhs` for the value of the matrix `A` that was last given to the constructor or to the `ComputeCholeskyDecomposition` function *at the time* it was passed to the function. The function fails with return value 101 if it is called after the `directIterative` option was enabled in the constructor; otherwise, returns the error status of the `pardiso` function.

```
int SolveLinearSystemDirectIterative(const SparseMatrix * A,
    double * x, const double * rhs)
```
Solves the linear system `Ax = rhs` with the direct iterative method, using the input matrix `A`. Returns the error status of the `pardiso` function.

class SPOOLESSolver : public LinearSolver Solves the linear system $Ax = b$ using the public domain SPOOLES solver. Matrix A must be symmetric and invertible.

```
SPOOLESSolver(const SparseMatrix * A, int verbose=0)
```
Initializes the solver with the symmetric matrix `A`, so that the value of `A` *at the time of construction* will be used for all calls of `SolveLinearSystem`. Setting `verbose` to 1 prints feedback during the constructor, and to 2 additionally prints feedback during the solve. Throws an exception upon failure.

```
virtual int SolveLinearSystem(double * x, const double * rhs)
```
Solves the system `Ax = rhs` using the matrix `A` from the constructor. Returns 0 iff successful.

class SPOOLESSolverMT : public LinearSolver   Solves the linear system $Ax = b$ using the public domain multithreaded SPOOLES solver. Matrix A must be symmetric and invertible.

    `SPOOLESSolverMT(const SparseMatrix * A, int numThreads, int verbose=0)`
Initializes the solver with the matrix `A`, so that the value of `A` *at the time of construction* will be used for all calls of `SolveLinearSystem`, and sets the number of threads `numThreads` to use during the solve. Throws an exception upon failure. Setting `verbose` to 1 prints feedback during the constructor, and to 2 additionally prints feedback during the solve. Further diagnostic information is written to a file called `SPOOLES.message`.

    `virtual int SolveLinearSystem(double * x, const double * rhs)`
Solves the linear system $Ax = \text{rhs}$ for the matrix `A` from the constructor, using `numThreads` threads as specified in the constructor. Returns the exit code of the solver (1 if successful).

2.46   stencilForceModel

Provides an interface that enables fast multi-threading across the elements (or stencils, to be precise) of a simulation mesh.

class StencilForceModel   Provides an interface to collect elastic energy, internal forces and stiffness matrix for each stencil in a mesh. A stencil is defined as the basic building block for computing simulation quantities on a mesh. For a volumetric mesh, each element is a stencil. For cloth simulation, each triangle is a stencil for in-plane stretch and shear energies, and the two triangles sharing an edge (a "bending unit") form a stencil for bending energy. For mass-spring simulation, each spring is a stencil.

    `virtual void GetStencilLocalEnergyAndForceAndMatrix( int stencilType, int stencilId,`
       `const double * u, double * energy, double * internalForces,`
       `double * tangentStiffnessMatrix ) = 0`
Computes the energy, internal forces, tangent stiffness matrix of a stencil with index `stencilId` and type `stencilType`. A volumetric mesh with m elements has one stencil type and this type has m stencils. A cloth mesh with m triangles and b bending units has two stencil types. The first type models in-place energy and has m stencils, and the second type models bending and has b stencils. A mass-spring system with m springs has one stencil type and this type has m stencils. Parameter `u` is the vertex displacements of the mesh. The output parameters `energy`, `internalForces` and `tangentStiffnessMatrix` return the computed energy, internal forces and dense tangentStiffnessMatrix for this stencil. If any of the three parameters are `nullptr`, the function will not compute the corresponding quantities.

class ForceModelAssembler   For each stencil types, assembles (adds) its values across the entire mesh. It inherits from `ForceModel` class and implements the necessary interfaces. If Intel TBB is provided, the computation will be performed in parallel, with the number of threads controlled by Intel TBB. The number of threads can be controlled outside the class using the Intel TBB APIs. If Intel TBB is not provided, the computation is single-threaded.

    `ForceModelAssembler(StencilForceModel * stencilForceModel)`
Creates a `ForceModelAssembler` to assemble stencils in `stencilForceModel`.

    `virtual void GetEnergyAndForceAndMatrix(const double * u, double * energy,`
       `double * internalForces, SparseMatrix * tangentStiffnessMatrix)`
Computes the energy, internal forces and tangent stiffness matrix of the mesh by assembling stencils. It is done in parallel if Intel TBB is enabled. Parameter `u` is the vertex displacements of the mesh, `energy` is the output energy of the mesh, and `internalForces` is a buffer with the same DOFs as the displacement (3 * #vertices) to hold the output internal forces. The output sparse matrix `tangentStiffnessMatrix` holds the gradient of the internal forces (the tangent stiffness matrix). If any of the three parameters are `nullptr`, the function will not compute the corresponding quantities.

```
void GetStiffnessMatrixTopology(SparseMatrix ** stiffnessMatrixTopology)
```
Writes to `*stiffnessMatrixTopology` a newly allocated (using `new`) matrix containing the pattern of non-zero entries of the stiffness matrix.

class ClothBWStencilForceModel : public StencilForceModel   Implements the cloth stencils using the common interface of `StencilForceModel`.

```
ClothBWStencilForceModel(ClothBW * clothBW)
```
Sets the `ClothBW` object for energy, force and stiffness calculations.

class CorotationalLinearFEMStencilForceModel : public StencilForceModel   Implements the corotational linear FEM stencils using the common interface of `StencilForceModel`.

```
CorotationalLinearFEMStencilForceModel(CorotationalLinearFEM *
    corotationalLinearFEM)
```
Sets the `CorotationalLinearFEM` object for energy, force and stiffness calculations.

```
void SetWarp(int warp)
```
Sets the warp parameter. This makes it possible to change the warp parameter at runtime.

class IsotropicHyperelasticFEMStencilForceModel : public StencilForceModel   Implements the isotropic hyperelastic FEM stencils using the common interface of `StencilForceModel`.

```
IsotropicHyperelasticFEMStencilForceModel(IsotropicHyperelasticFEM *
    isotropicHyperelasticFEM )
```
Sets the `IsotropicHyperelasticFEM` object for energy, force and stiffness calculations.

class MassSpringStencilForceModel : public StencilForceModel   Implements the mass spring system stencils using the common interface of `StencilForceModel`.

```
MassSpringStencilForceModel(MassSpringSystem * system)
```
Sets the `MassSpringSystem` object for energy, force and stiffness calculations.

class StVKStencilForceModel : public StencilForceModel   Implements the StVK stencils using the common interface of `StencilForceModel`.

```
StVKStencilForceModel(StVKFEM * stvkFEM )
```
Sets the `StVKFEM` object for energy, force and stiffness calculations.

class LinearFEMStencilForceModel : public StencilForceModel   Creates a force model to linearize any deformable models.

```
LinearFEMStencilForceModel(StencilForceModel * stencilForceModel)
```
Linearizes `stencilForceModel`, calculating energy, force and stiffness as if `stencilForceModel` is a linear deformable model.

2.47   stvk

Implements the St.Venant-Kirchhoff (StVK) deformable model, providing classes to evaluate energy, internal forces, stiffness matrix, and Hessian tensor for a given volumetric mesh and vertex displacements.

class StVKElementABCD   Serves as a base class for classes that generate the St. Venant-Kirchhoff ABCD coefficients (see [Bar07]) for tetrahedral or cubic mesh elements.

```
virtual void AllocateElementIterator(void ** elementIterator)
```
Allocates and returns a pointer to a data structure used to hold information about a single mesh element for use in the ABCD calculations. The base class allocates a dummy structure, but implementations allocate structures of varying types, so `void` pointers are necessary.

```
virtual void ReleaseElementIterator(void * elementIterator)
```
De-allocates the data structure produced by `AllocateElementIterator`.

```
virtual void PrepareElement(int el, void * elementIterator)
```
Writes to the pre-allocated iterator `elementIterator` the appropriate pre-calculated information for the element at index `el` in the volumetric mesh.

```
virtual Mat3d A(void * elementIterator, int i, int j) = 0
virtual double B(void * elementIterator, int i, int j) = 0
virtual Vec3d C(void * elementIterator, int i, int j, int k) = 0
virtual double D(void * elementIterator, int i, int j, int k, int l) = 0
```
Calculate the A, B, C, or D values for the element whose iterator `elementIterator` has been allocated and prepared by the functions above, and for the chosen vertices `i`, `j`, `k`, and `l` of the element. Each integer ranges from 0 to (number of vertices per element)−1.

class StVKCubeABCD : public StVKElementABCD   Implements the base ABCD coefficients class for cube-shaped elements.

```
StVKCubeABCD(double cubeSize)
```
Initializes the class given the undeformed edge length `cubeSize` of the cubes in the cubic mesh. No other information about the mesh is needed to be able to compute the ABCD values, which are identical for each element in the mesh, though the ABCD functions still accept the `elementIterator` argument for consistency with the base class.

class StVKTetABCD : public StVKElementABCD   Implements the base ABCD coefficients class for tetrahedral elements.

```
StVKTetABCD(TetMesh * tetMesh)
```
Reads the mesh `tetMesh` to extract element-specific information used during the ABCD calculations. This information is distributed to the ABCD calculation functions via the `elementIterator` object produced by the class.

class StVKTetHighMemoryABCD : public StVKElementABCD   Implements the base ABCD coefficients class for tetrahedral elements, optimizing ABCD access speed at the expense of greater memory usage, by explicitly precalculating all ABCD values at the time of construction.

```
StVKTetHighMemoryABCD(TetMesh * tetMesh)
```
Reads the mesh `tetMesh` to extract element-specific information, and uses this information to calculate all the ABCD coefficients for each element. As a result, the ABCD calculation function implementations of this subclass consist of a single array lookup.

class StVKElementABCDLoader   Provides a single function to create an appropriate `StVKElementABCD` object for a given volumetric mesh.

```
static StVKElementABCD * load(VolumetricMesh * volumetricMesh,
    unsigned int loadingFlag=0)
```
Given the mesh `volumetricMesh`, allocates (`new`) and returns the pointer to an instance of a child class of `StVKElementABCD` which can produce the appropriate ABCD values for the mesh. Setting `loadingFlag` to 1 will instruct the loader to create `StVKTetHighMemoryABCD` instances for tet meshes. Returns `NULL` or calls `exit` with non-zero status upon failure.

class StVKFEM   Calculates elastic energy, internal forces and tangent stiffness matrix for the StVK material.

```
StVKFEM(VolumetricMesh * volumetricMesh, StVKElementABCD * precomputedABCDIntegrals)
```
Initializes the class to use mesh `volumetricMesh` and the precalculated information `precomputedABCD-Integrals` for that mesh (usually generated by `StVKElementABCDLoader` given the mesh).

```
void ComputeElementLocalEnergyAndInternalForcesAndStiffnessMatrix(
    const double * vertexDisplacements, int elementID, double * energy,
    double * internalForces, double * stiffnessMatrix)
```
Computes the non-linear elastic strain energy, internal forces and stiffness matrix on one volumetric mesh element `elementID`, given a vector `vertexDisplacements` of the displacements for the vertices of the volumetric mesh. If `energy`, `internalForces` or `stiffnessMatrix` is NULL, the function does not calculate or return the corresponding information.

2.48   virtualTets

Implements our "virtual tets" method for nearly self-intersecting triangle meshes. Creates a virtual tet mesh that embeds an intersection-free triangle mesh (that may be nearly self-intersecting), given an input tet mesh that covers the space occupied by the triangle mesh. The triangle mesh is assumed to be manifold, free of self-intersections and have no degenerate triangles. The method duplicates the tets and tet vertices to follow the topology of the triangle mesh. The implementation follows this paper:

```
Yijing Li, Jernej Barbic:
Immersion of Self-Intersecting Solids and Surfaces,
ACM Transactions on Graphics 37(4) (SIGGRAPH 2018)
```

```
TetMeshGeo createVirtualTetsMesh( const TetMeshGeo & tetMesh, const TriMeshGeo & triMesh,
    BarycentricCoordinates * coord = nullptr,
    std::vector<int> * newTetVtxID2OldTetVtxID = nullptr,
    std::vector<int> * newTetID2OldTetID = nullptr,
    std::vector<std::vector<int>> * tetTris = nullptr, bool verbose = false,
    Profiler * profiler = nullptr)
```
Input `triMesh` and `tetMesh` are the triangle and tet mesh geometry. The tet mesh must cover the space occupied by the triangle mesh. The returned `TetMeshGeo` is the tet mesh where the tets and tet vertices are duplicated to follow the topology of the triangle mesh. Optionally, the routine returns `coord` which stores the barycentric weights to embed the triangle mesh into the returned tet mesh, the buffer `newTetVtxID2OldTetVtxID` as a mapping from new tet vertex indices to input tet vertex indices, `newTetID2OldTetID` as a mapping from new tet indices to input tet indices, and `tetTris` as a mapping from each new tet index to the triangle indices embedded in this tet. Parameter `verbose` controls verbosity, and the optional `profiler` collects performance data.

2.49   volumetricMesh

Provides tools to load and store tetrahedral or cubic volumetric meshes, set and access their material properties, and perform various calculations with them.

class VolumetricMesh Serves as an abstract base class for volumetric meshes, holding the vertices and elements of the mesh, identifying regions of the mesh with common material properties and listing those properties, and providing functionality common to meshes with different element geometries. Derived into `TetMesh` and `CubicMesh` to store tet meshes and meshes consisting of cubes (voxels). `VolumetricMesh` is abstract; you cannot initialize it directly, but must be initialized via one of its two subclasses. Note: The loading procedure is made easier by the `VolumetricMeshLoader` class.

```
static elementType getElementType(const char * filename)
```

```
virtual elementType getElementType() = 0
```
Returns the element type of the mesh in file `filename` or of the current mesh. The result is "TET" for `TetMesh` objects and "CUBIC" for `CubicMesh` objects. If `filename` cannot be read, the type is set to `INVALID`.

```
int getNumElementVertices()
```
Returns the number of vertices per mesh element.

```
virtual int saveToAscii(const char * filename) = 0
```
Saves the mesh, including materials, sets and regions, to a `filename`. Format is `.veg`, which may be later loaded by `VolumetricMeshLoader`.

```
virtual int saveToBinary(const char * filename, unsigned int * bytesWritten) = 0
```
Saves the mesh, including materials, sets and regions, to a `filename`, in binary format. If a non-NULL `bytesWritten` is provided, the function will return in `*bytesWritten` the number of bytes written.

```
int getNumVertices()
int getNumElements()
int getNumMaterials()
int getNumSets()
int getNumRegions()
```
Returns the number of various mesh objects.

```
Vec3d & getVertex(int i)
Material * getMaterial(int i)
Set * getSet(int i)
Region * getRegion(int i)
```
Returns a reference or pointer to the specified type of object, given its index in the list of that type of object in the `VolumetricMesh`.

```
void setMaterial(int i, const Material * material)
```
Sets the material at index `i` in the array of materials to `material`, and copies these material properties to each mesh element whose containing region lists index `i` as the material which defines its material properties.

```
virtual double getElementVolume(int el) = 0
```
Implementations return the volume of the element at index `el`.

```
virtual bool containsVertex(int element, Vec3d pos) = 0
```
Returns whether the element at index `i` contains the vector `pos`.

```
int getClosestElement(Vec3d pos)
int getClosestVertex(Vec3d pos)
int getContainingElement(Vec3d pos)
```
Returns the index of the nearest element, nearest vertex, or containing element of `pos` in the undeformed mesh. For the third function, -1 is returned if no containing element is found.

```
void exportMeshGeometry(int * numVertices, double ** vertices,
    int * numElements, int * numElementVertices, int ** elements)
```
Exports the geometric information of the mesh to memory arrays. The number of vertices, number of elements, and number of vertices per element are written to `*numVertices`, `*numElements`, and `*numElementVertices`. Array `*vertices` is set to point to a newly-allocated (`malloc`) array of the coordinates of all the vectors, and `*elements` is set to point to a newly-allocated (`malloc`) array of the indices of the vertices constituting each mesh element.

```
int generateInterpolationWeights(int numTargetLocations,
    double * targetLocations, int ** vertices, double ** weights,
    double zeroThreshold = -1.0, int ** elements = NULL,
    int verbose=0)
```
Generates interpolation information for embedding a (typically higher-resolution) rendering mesh within a deformable volumetric mesh. Given `numTargetLocations` points (rendering mesh vertices), with coordinates specified in the array `targetLocations`, the function determines which mesh element is closest to or contains each point. It then expresses each point's position as a weighted average of the coordinates of that element's vertices, such that the weights sum to one (standard tet barycentric coordinates). Note that weights may be negative if the point is outside the volumetric mesh. Output arrays `*vertices` and `*weights` are allocated (`malloc`) inside the function. For each target location, the `*vertices` array gives the integer indices of `VolumetricMesh` vertices of the mesh element containing the target. The array `*weights` gives the barycentric weights. The length of both arrays is `numElementVertices` · `numTargetLocations`. If `zeroThreshold` > 0, then the points that are further than `zeroThreshold` away from any volumetric mesh vertex, are assigned weights of 0. If `elements` is not NULL, the function will allocate an integer array `*elements`, and return the closest element to each target location in it. The function returns the number of target points which do not lie inside any element.

```
static int loadInterpolationWeights(const char * filename,
    int numTargetLocations, int numElementVertices,
    int ** vertices, double ** weights)
static int saveInterpolationWeights(const char * filename,
    int numTargetLocations, int numElementVertices,
    int * vertices, double * weights)
```
Reads or writes the arrays `vertices` and `weights` of interpolation information from or to `filename`. Returns 0 iff successful.

```
static void interpolate(double * u, double * uTarget, int numTargetLocations,
    int numElementVertices, int * vertices, double * weights)
```
Interpolates volumetric mesh vertex deformations `u` to an embedded (triangle, volumetric, etc.) mesh, using barycentric linear interpolation. The computed embedded mesh vertex deformations are stored into the pre-allocated array `uTarget`. The number of the embedded mesh vertices must be provided in `numTargetLocations`. The number of vertices in an element of the volumetric mesh must be provided in `numElementVertices`. Input parameters `vertices` and `weights` must have been generated by a previous call to `generateInterpolationWeights`. Note that `u` can be set to either the deformed positions or the displacements of the volumetric mesh vertices to obtain in `uTarget` either the deformed positions or the displacements of the target points.

```
VolumetricMesh(int numVertices, double * vertices,
    int numElements, int numElementVertices, int * elements,
    double E=1E6, double nu=0.45, double density=1000)
```
This non-public constructor is called by derived classes' constructors. It constructs a mesh with `numVertices` vertices whose coordinates are given in the size 3 · `numVertices` array `vertices` and with `numElements` elements (each with `numElementVertices` vertices), whose constituent vertices are specified by the `numElements`· `numElementVertices` size array of vertex indices `elements`.

class VolumetricMesh::Material   A nested subclass of `VolumetricMesh`. Stores material properties. Constructed by the `VolumetricMesh` constructor. This class is abstract and stores common functionality for storing materials. Derived into classes `ENuMaterial` and `MooneyRivlinMaterial`.

class VolumetricMesh::ENuMaterial   The `ENuMaterial` class stores material properties for any material that is parameterized by Young's modulus (E) and Poisson's ratio (ν). Most materials in Vega falls into this category, e.g., co-rotational linear FEM, StVK, invertible StVK, invertible neo-Hookean, etc. Constructed by the `VolumetricMesh` constructor.

```
double getE()
double getNu()
double getDensity()
double getLambda()
double getMu()
```
Return the `E`, `nu`, or `density` property of the material, as well as the Lamé coefficients λ and μ (computed from E, ν).

```
std::string getName()
```
Returns the name string of the material.

```
void setE(double E)
void setNu(double nu)
void setDensity(double density)
void setName(char * name)
```
Set the relevant material property or the name of the material.

class VolumetricMesh::MooneyRivlinMaterial   Stores a Mooney-Rivlin material, parameterized by μ_{01}, μ_{10}, v_1. Sister class to `ENuMaterial`. Constructed by the `VolumetricMesh` constructor.

class VolumetricMesh::Set   A nested subclass of `VolumetricMesh`. Stores a `set` of integer indices. Constructed by the `VolumetricMesh` constructor. Typically, it is used to store indices of elements that use the same material.

```
void addElement(int element)
```
Adds `element` to the set.

```
void getElements(std::set<int> & elements)
```
Stores a copy of the set elements to the set `elements`.

class VolumetricMesh::Region   A region is a subpart of the mesh consisting of common material properties. A nested subclass of `VolumetricMesh`. Constructed by the `VolumetricMesh` constructor.

```
int getMaterialIndex()
```
Returns the index of the `VolumetricMesh` material which corresponds to the mesh elements in this region.

```
int getSetIndex()
```
Returns the index of the `VolumetricMesh` set which corresponds to the mesh elements in this region.

class CubicMesh : public VolumetricMesh   Stores a cubic mesh in which each mesh element is a cube. All cubes have the same size, and must originate from an axis-aligned grid of uniform size. Such meshes can be easily created from input triangular geometry using the LargeModalDeformationFactory utility.

```
CubicMesh(int numVertices, double * vertices, int numElements,
    int * elements, double E=1E6, double nu=0.45, double density=1000)
```
Constructs a cubic mesh of `numVertices` vertices with coordinates in the array `vertices`, organized into `numElements` cubes by the array `elements` which contains eight vertex indices for each cube. Assumes that the input actually specifies cube-shaped elements that come from an axis-aligned grid. Upon failure, throws an exception.

```
CubicMesh(const char * filename, fileFormatType fileFormat=ASCII, int verbose=1)
```
Loads the mesh from the file `filename`. If format is ASCII, the file is in the text `.veg` format. If format is BINARY, the file is in the binary `.vegb` format. Throws an exception on failure.

```
static CubicMesh * CreateFromUniformGrid(int resolution,
    int numVoxels, int * voxels, double E=1E6, double nu=0.45, double density=1000)
```
Constructs a mesh that is a subset of a uniform voxel grid of dimensions resolution^3. The `numVoxels` mesh elements (voxels) are specified by their integer coordinates in the integer array `voxels` (length 3·`numVoxels`). The function returns a pointer to a newly-allocated (`new`) `CubicMesh` with this geometry.

```
double cubeSize()
```
Returns the grid size (=length of an edge) of the cubes in the mesh.

class TetMesh : public VolumetricMesh   Stores a tetrahedral mesh. The tetrahedral mesh elements may be arbitrarily shaped. To create tetrahedral meshes, you can use a mesh generation package such as TetGen [Han11] or Stellar [KS09, Kli08].

```
TetMesh(const char * filename, fileFormatType fileFormat=ASCII, int verbose=1)
```
Loads the mesh from the file `filename`. If format is ASCII, the file is in the text `.veg` format. If format is BINARY, the file is in the binary `.vegb` format. Throws an exception on failure.

```
void orient()
```
Orients all the tets of the mesh (re-orders vertices within each tet), so that each tet has positive orientation: $((v_1 - v_0) \times (v_2 - v_0)) \cdot (v_3 - v_0) \geq 0$. If your input .veg mesh does not have this property, you should call `orient` after loading the mesh. Note: such positive orientation is required for several of the Vega material models to function properly. Note: meshes produced by TetGen already are oriented positively, so there is no need to call `orient` for meshes produced by TetGen.

class VolumetricMeshLoader   Provides a single function to load cubic or tetrahedral meshes. It is not necessary to know the mesh element type in advance.

`static VolumetricMesh * load(const char * filename)`   Loads the mesh in file `filename` and returns the pointer to a newly-allocated (`new`) `CubicMesh` instance or `TetMesh`. Returns `NULL` or throws an exception upon failure, depending on which subfunction experiences the failure.

Acknowledgments

References

[Bar07]   Jernej Barbič. *Real-time Reduced Large-Deformation Models and Distributed Contact for Computer Graphics and Haptics.* PhD thesis, Carnegie Mellon University, August 2007.

[Bar12]   Jernej Barbič. Exact Corotational Linear FEM Stiffness Matrix. Technical report, University of Southern California, 2012.

[BJ05]   Jernej Barbič and Doug L. James. Real-time subspace integration for St. Venant-Kirchhoff deformable models. *ACM Trans. on Graphics*, 24(3):982–990, 2005.

[Bow81]   A. Bowyer. Computing dirichlet tessellations. *Computer Journal*, 24(2):162–166, 1981.

[Bow09]   Allan Bower. *Applied Mechanics of Solids.* CRC Press, 2009.

[BW98]   David Baraff and Andrew P. Witkin. Large Steps in Cloth Simulation. In *Proc. of ACM SIGGRAPH 98*, pages 43–54, July 1998.

[BW01]   David Baraff and Andrew Witkin. Physically based modelling, ACM SIGGRAPH Course Notes, 2001.

[BW08]   Javier Bonet and Richard D. Wood. *Nonlinear Continuum Mechanics for Finite Element Analysis, 2nd Ed.* Cambridge University Press, 2008.

[Che95]   E. V. Chernyaev. Marching Cubes 33: construction of topologically correct isosurfaces. Technical report, Technical Report CERN CN 9517, 1995.

[CPSS10]   Isaac Chao, Ulrich Pinkall, Patrick Sanan, and Peter Schröder. A Simple Geometric Model for Elastic Deformations. *ACM Transactions on Graphics*, 29(3):38:1–38:6, 2010.

[Han11]   Hang Si. TetGen: A Quality Tetrahedral Mesh Generator and a 3D Delaunay Triangulator, 2011. http://tetgen.org/.

[HX16]   Jernej Barbič Hongyi Xu. Pose-space subspace dynamics. *ACM Trans. on Graphics (SIGGRAPH 2016)*, 35(4), 2016.

[ITF04]   G. Irving, J. Teran, and R. Fedkiw. Invertible Finite Elements for Robust Simulation of Large Deformation. In *Symp. on Computer Animation (SCA)*, pages 131–140, 2004.

[JKSH13]   Alec Jacobson, Ladislav Kavan, , and Olga Sorkine-Hornung. Robust inside-outside segmentation using generalized winding numbers. *ACM Transactions on Graphics*, 32(4):33:1–33:12, 2013.

[Kli08]   Matthew Klingner. *Improving Tetrahedral Meshes.* PhD thesis, Department of Electrical Engineering and Computer Sciences, University of California, Berkeley, 2008.

[KS09]   Matthew Klingner and Jonathan Richard Shewchuk. Stellar: A tetrahedral mesh improvement program, 2009. http://www.cs.berkeley.edu/~jrs/stellar/.

[KTY09]   Ryo Kikuuwe, Hiroaki Tabuchi, and Motoji Yamamoto. An edge-based computationally efficient formulation of saint venant-kirchhoff tetrahedral finite elements. *ACM Trans. on Graphics*, 28(1):1–13, 2009.

[LB14]   Yijing Li and Jernej Barbič. Stable orthotropic materials. In *Symp. on Computer Animation (SCA)*, 2014.

[LLVT03]   T. Lewiner, H. Lopes, A. W. Vieira, and G. Tavares. Efficient implementation of Marching Cubes' cases with topological guarantees. *Journal of Graphics Tools*, 8(2):1–15, 2003.

[MG04]   M. Müller and M. Gross. Interactive Virtual Materials. In *Proc. of Graphics Interface 2004*, pages 239–246, 2004.

[ORY05] Steve Oudot, Laurent Rineau, and Mariette Yvinec. Meshing volumes bounded by smooth surfaces. In *Proc. of 14th International Meshing Roundtable*, pages 203–220, 2005.

[PAR] PARDISO: Parallel Direct Sparse Solver Interface. Pardiso project, www.pardiso-project.org and Intel MKL, software.intel.com/en-us/articles/intel-mkl.

[SA07] Olga Sorkine and Marc Alexa. As-rigid-as-possible surface modeling. In *Symp. on Geometry processing*, volume 4, pages 109–116, 2007.

[SG11] Hang Si and Klaus Gärtner. 3d boundary recovery by constrained delaunay tetrahedralization. *International Journal for Numerical Methods in Engineering*, 85(11):1341–1364, 2011.

[Sha90] Ahmed A. Shabana. *Theory of Vibration, Volume II: Discrete and Continuous Systems*. Springer–Verlag, 1990.

[She94] Jonathan Richard Shewchuk. An introduction to the conjugate gradient method without the agonizing pain, 1994.

[She97] Jonathan Richard Shewchuk. Adaptive precision floating-point arithmetic and fast robust geometric predicates. *Discrete & Computational Geometry*, 18(3):305–363, 1997.

[Si08] Hang Si. Adaptive tetrahedral mesh generation by constrained Delaunay refinement. *International Journal for Numerical Methods in Engineering*, 75:856–880, 2008.

[Sid98] Roger B. Sidje. Expokit: A Software Package for Computing Matrix Exponentials. *ACM Trans. on Mathematical Software*, 24(1):130–156, 1998. www.expokit.org.

[SPO] SPOOLES: SParse Object Oriented Linear Equations Solver. Boeing Phantom Works. www.netlib.org/linalg/spooles/spooles.2.2.html.

[SSB12] Fun Shing Sin, Daniel Schroeder, and Jernej Barbič. Vega: Nonlinear FEM Deformable Object Simulator. *Computer Graphics Forum*, 2012. accepted, to appear.

[TSIF05] Joseph Teran, Eftychios Sifakis, Geoffrey Irving, and Ronald Fedkiw. Robust Quasistatic Finite Elements and Flesh Simulation. In *Symp. on Computer Animation (SCA)*, pages 181–190, 2005.

[Wri02] Peter Wriggers. *Computational Contact Mechanics*. John Wiley & Sons, Ltd., 2002.

[XB14] Hongyi Xu and Jernej Barbič. Signed distance fields for polygon soup meshes. In *Proc. of Graphics Interface*, pages 35–41, 2014.

[XLCB15] Hongyi Xu, Yijing Li, Yong Chen, and Jernej Barbič. Interactive material design using model reduction. *ACM Trans. on Graphics*, 34(2), 2015.